

project2

May 15, 2025

Project - Predictive Analysis for Credit Card Fraud Detection Using Classification Models

Problem Statement:

In the financial industry, fraudulent credit card transactions can cause significant losses and damage customer trust. The challenge lies in identifying fraudulent transactions among millions of legitimate ones, especially when fraud cases represent less than 0.2% of all transactions. This project aims to develop and evaluate classification models that can accurately detect credit card fraud, even in the presence of severe class imbalance.

Dataset Information:

Source: Kaggle - Credit Card Fraud Detection Dataset

Link: [Kaggle Dataset](#)

Records: 284,807 transactions

Features: 30 columns (28 PCA features V1-V28, plus Time, Amount)

Target Column: Class (0 = legitimate, 1 = fraud)

Initial Exploration & Preprocessing:

1 Data Loading from ZIP:

Loaded creditcard.csv directly from ZIP using `zipfile.ZipFile`

2 Exploratory Data Analysis (EDA):

Used `df.describe()` to understand feature distributions

Observed heavy class imbalance (0 » 1)

Univariate Analysis:

Plotted histograms for individual feature columns using `sns.histplot`

Identified skewness in many features

Bivariate Analysis:

Plotted distributions of each feature vs. Class using overlaid histograms

Helped identify which features separate frauds from normal transactions

Target Analysis:

Fraud class (1) is only 0.17% of data

3 Data Preprocessing:

Feature Scaling:

Scaled Amount and Time using StandardScaler

Splitting Data:

Used `train_test_split(X, y, test_size=0.2, stratify=y)` to ensure balanced splits

Handling Class Imbalance:

Used SMOTE to oversample minority class in training data

Post-SMOTE: Balanced dataset with equal 0s and 1s

4 Model Building:

Algorithm used to create a model:

Logistic Regression (Baseline)

Random Forest Classifier

K-Nearest Neighbors (KNN)

Loading libraries and Data

```
[2]: # import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import zipfile
# Suppress future warnings
import warnings
warnings.filterwarnings('ignore')

[4]: # Define the path to the zip file dataset
zip_path = r'C:\Users\hp\Downloads\creditcard.csv.zip'

# Open the zip file
with zipfile.ZipFile(zip_path) as z:

    # Get the CSV filename (first file in the archive)
    csv_file = z.namelist()[0]

    # Read the CSV file
    with z.open(csv_file) as f:
```

```
# Read the CSV file
data = pd.read_csv(f)
```

```
# Display the first 5 rows
print(data.head())
```

	Time	V1	V2	V3	V4	V5	V6	V7	\
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	

	V8	V9	...	V21	V22	V23	V24	V25	\
0	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928	0.128539	
1	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846	0.167170	
2	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689281	-0.327642	
3	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575	0.647376	
4	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0.141267	-0.206010	

	V26	V27	V28	Amount	Class
0	-0.189115	0.133558	-0.021053	149.62	0
1	0.125895	-0.008983	0.014724	2.69	0
2	-0.139097	-0.055353	-0.059752	378.66	0
3	-0.221929	0.062723	0.061458	123.50	0
4	0.502292	0.219422	0.215153	69.99	0

[5 rows x 31 columns]

```
[5]: # Returns a tuple representing (rows, columns)
data.shape
```

[5]: (284807, 31)

concept of pca , what are principle componenets , aurthonalaty , multyquonarity , eigen values , eigen vector

Column Dictionary

Time - Time in seconds elapsed between this transaction and the first transaction in the dataset.

V1-V28 - Principal components obtained from a PCA (Principal Component Analysis) transformation to anonymize the features (original features like name, merchant, location are not available).

Amount - The transaction amount. Useful for scaling and normalization.

Class - Target variable: 0 for non-fraud, 1 for fraud.

```
[7]: # see column data type and some info
data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Time        284807 non-null  float64
1   V1          284807 non-null  float64
2   V2          284807 non-null  float64
3   V3          284807 non-null  float64
4   V4          284807 non-null  float64
5   V5          284807 non-null  float64
6   V6          284807 non-null  float64
7   V7          284807 non-null  float64
8   V8          284807 non-null  float64
9   V9          284807 non-null  float64
10  V10         284807 non-null  float64
11  V11         284807 non-null  float64
12  V12         284807 non-null  float64
13  V13         284807 non-null  float64
14  V14         284807 non-null  float64
15  V15         284807 non-null  float64
16  V16         284807 non-null  float64
17  V17         284807 non-null  float64
18  V18         284807 non-null  float64
19  V19         284807 non-null  float64
20  V20         284807 non-null  float64
21  V21         284807 non-null  float64
22  V22         284807 non-null  float64
23  V23         284807 non-null  float64
24  V24         284807 non-null  float64
25  V25         284807 non-null  float64
26  V26         284807 non-null  float64
27  V27         284807 non-null  float64
28  V28         284807 non-null  float64
29  Amount      284807 non-null  float64
30  Class       284807 non-null  int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB

```

```

[9]: # see missing value in each column
data.isna().sum()

```

```

[9]: Time      0
     V1        0
     V2        0
     V3        0
     V4        0

```

```

V5      0
V6      0
V7      0
V8      0
V9      0
V10     0
V11     0
V12     0
V13     0
V14     0
V15     0
V16     0
V17     0
V18     0
V19     0
V20     0
V21     0
V22     0
V23     0
V24     0
V25     0
V26     0
V27     0
V28     0
Amount  0
Class   0
dtype: int64

```

```
[19]: data.describe()
```

```

[19]:

```

	Time	V1	V2	V3	V4 \
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	94813.859575	1.168375e-15	3.416908e-16	-1.379537e-15	2.074095e-15
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01

	V5	V6	V7	V8	V9 \
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	9.604066e-16	1.487313e-15	-5.556467e-16	1.213481e-16	-2.406331e-15
std	1.380247e+00	1.332271e+00	1.237094e+00	1.194353e+00	1.098632e+00
min	-1.137433e+02	-2.616051e+01	-4.355724e+01	-7.321672e+01	-1.343407e+01
25%	-6.915971e-01	-7.682956e-01	-5.540759e-01	-2.086297e-01	-6.430976e-01
50%	-5.433583e-02	-2.741871e-01	4.010308e-02	2.235804e-02	-5.142873e-02

75%	6.119264e-01	3.985649e-01	5.704361e-01	3.273459e-01	5.971390e-01
max	3.480167e+01	7.330163e+01	1.205895e+02	2.000721e+01	1.559499e+01

	...	V21	V22	V23	V24 \
count	...	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	...	1.654067e-16	-3.568593e-16	2.578648e-16	4.473266e-15
std	...	7.345240e-01	7.257016e-01	6.244603e-01	6.056471e-01
min	...	-3.483038e+01	-1.093314e+01	-4.480774e+01	-2.836627e+00
25%	...	-2.283949e-01	-5.423504e-01	-1.618463e-01	-3.545861e-01
50%	...	-2.945017e-02	6.781943e-03	-1.119293e-02	4.097606e-02
75%	...	1.863772e-01	5.285536e-01	1.476421e-01	4.395266e-01
max	...	2.720284e+01	1.050309e+01	2.252841e+01	4.584549e+00

	V25	V26	V27	V28	Amount \
count	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	284807.000000
mean	5.340915e-16	1.683437e-15	-3.660091e-16	-1.227390e-16	88.349619
std	5.212781e-01	4.822270e-01	4.036325e-01	3.300833e-01	250.120109
min	-1.029540e+01	-2.604551e+00	-2.256568e+01	-1.543008e+01	0.000000
25%	-3.171451e-01	-3.269839e-01	-7.083953e-02	-5.295979e-02	5.600000
50%	1.659350e-02	-5.213911e-02	1.342146e-03	1.124383e-02	22.000000
75%	3.507156e-01	2.409522e-01	9.104512e-02	7.827995e-02	77.165000
max	7.519589e+00	3.517346e+00	3.161220e+01	3.384781e+01	25691.160000

	Class
count	284807.000000
mean	0.001727
std	0.041527
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

[8 rows x 31 columns]

Univariate Analysis & Visualizations

```
[10]: #separate the features and target value
A = data.drop("Class", axis=1) # Features
B = data["Class"]             # Target
```

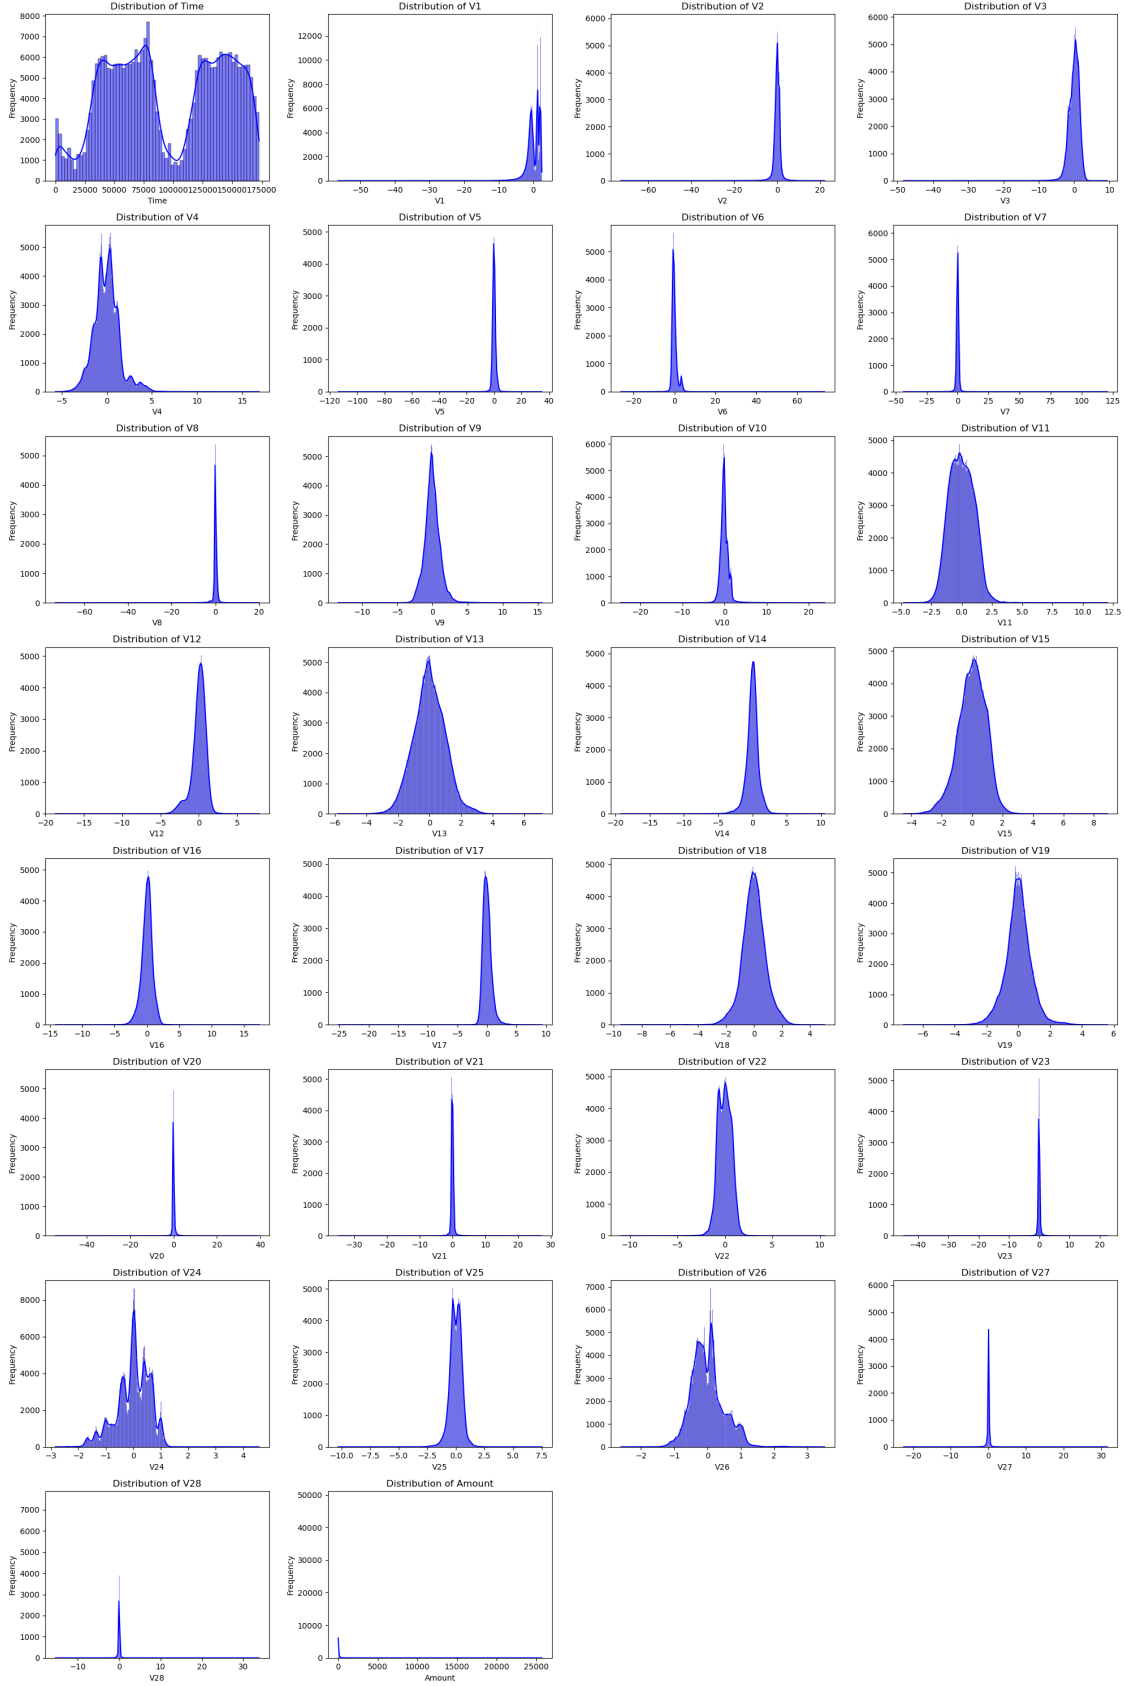
```
[35]: #Visualizations of Features value
# Create a histplot for the feature variable

plt.figure(figsize=(20, 30))

for i , col in enumerate(A):
    plt.subplot(8, 4, i + 1)
```

```
sns.histplot(data=data, x=col, kde=True, color='blue')
plt.title(f'Distribution of {col}')
plt.xlabel(col)
plt.ylabel('Frequency')

plt.tight_layout()
plt.show()
```



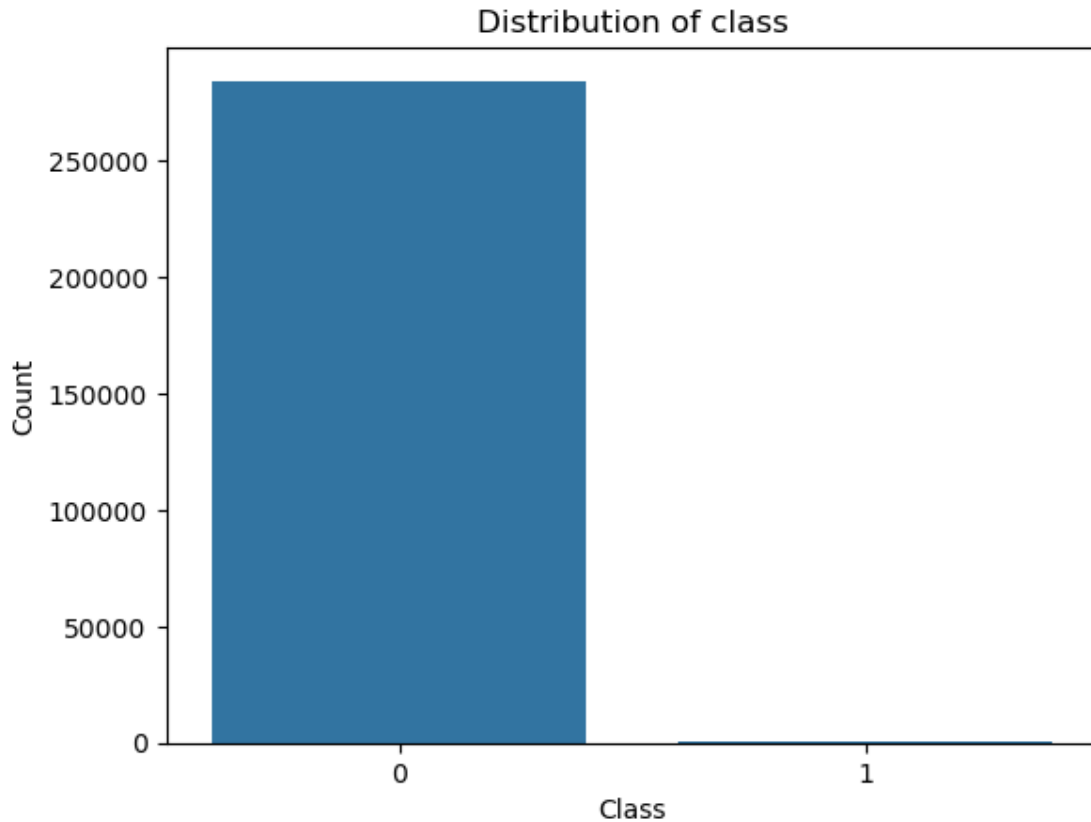
Columns Distribution

Time - Insight : The distribution shows a right skew, indicating that most transactions occur within a certain time frame, with fewer transactions occurring at the extremes. This suggests that the majority of transactions are clustered around specific times, while a few occur at much later times.

V1-V28 - Insight : These features exhibit various distributions. Some columns show a more concentrated range of values, while others have a wider spread. The presence of both positive and negative values indicates that these features capture different aspects of the transactions, possibly related to the underlying patterns of the data.

Amount - Insight : The histogram indicates that most transactions are relatively small, with a few outliers representing significantly larger amounts. This is typical in financial datasets where most transactions are low-value, but there are occasional high-value transactions that could be significant.

```
[11]: #Visualizations of Target value  
# Create a count plot for the target variable 'Class'  
sns.countplot(x=B, data=data)  
plt.title(f'Distribution of {"class"}')  
plt.xlabel("Class")  
plt.ylabel('Count')  
plt.show()
```



Class - Insight : The count plot shows a significant imbalance between the classes, with a vast majority of transactions being non- fraudulent (0) compared to fraudulent (1). This imbalance is crucial for modeling, as it may require techniques to handle class imbalance effectively.

Bivariate Analysis & Visualizations

```
[15]: # Calculate correlations between all features and the 'Class' column
correlations = data.corr()["Class"].sort_values(ascending=False)
print(correlations)
```

```
Class      1.000000
V11        0.154876
V4          0.133447
V2          0.091289
V21         0.040413
V19         0.034783
V20         0.020090
V8          0.019875
V27         0.017580
V28         0.009536
Amount     0.005632
V26         0.004455
```

```
V25      0.003308
V22      0.000805
V23     -0.002685
V15     -0.004223
V13     -0.004570
V24     -0.007221
Time    -0.012323
V6      -0.043643
V5      -0.094974
V9      -0.097733
V1      -0.101347
V18     -0.111485
V7      -0.187257
V3      -0.192961
V16     -0.196539
V10     -0.216883
V12     -0.260593
V14     -0.302544
V17     -0.326481
Name: Class, dtype: float64
```

Observations

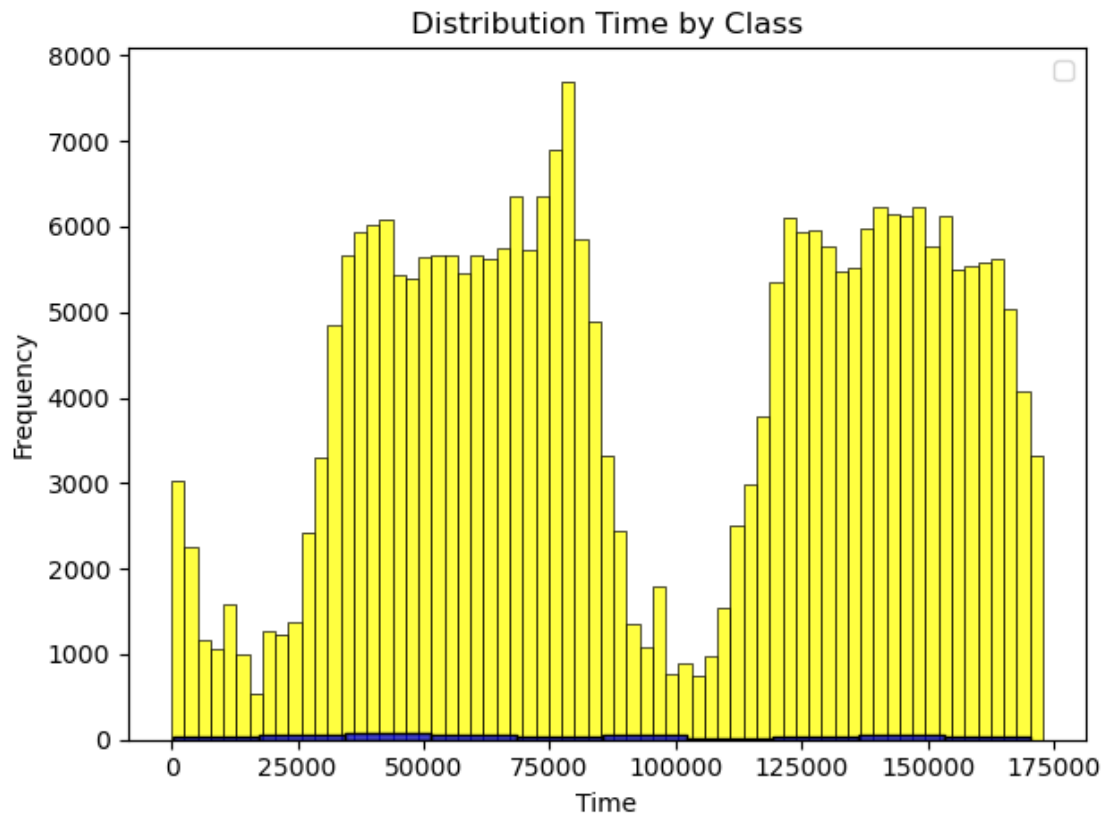
Strong Correlations: The features V11 and V4 show a strong positive correlation, indicating that as one increases, the other tends to increase as well. This could suggest that they are capturing similar underlying patterns in the data.

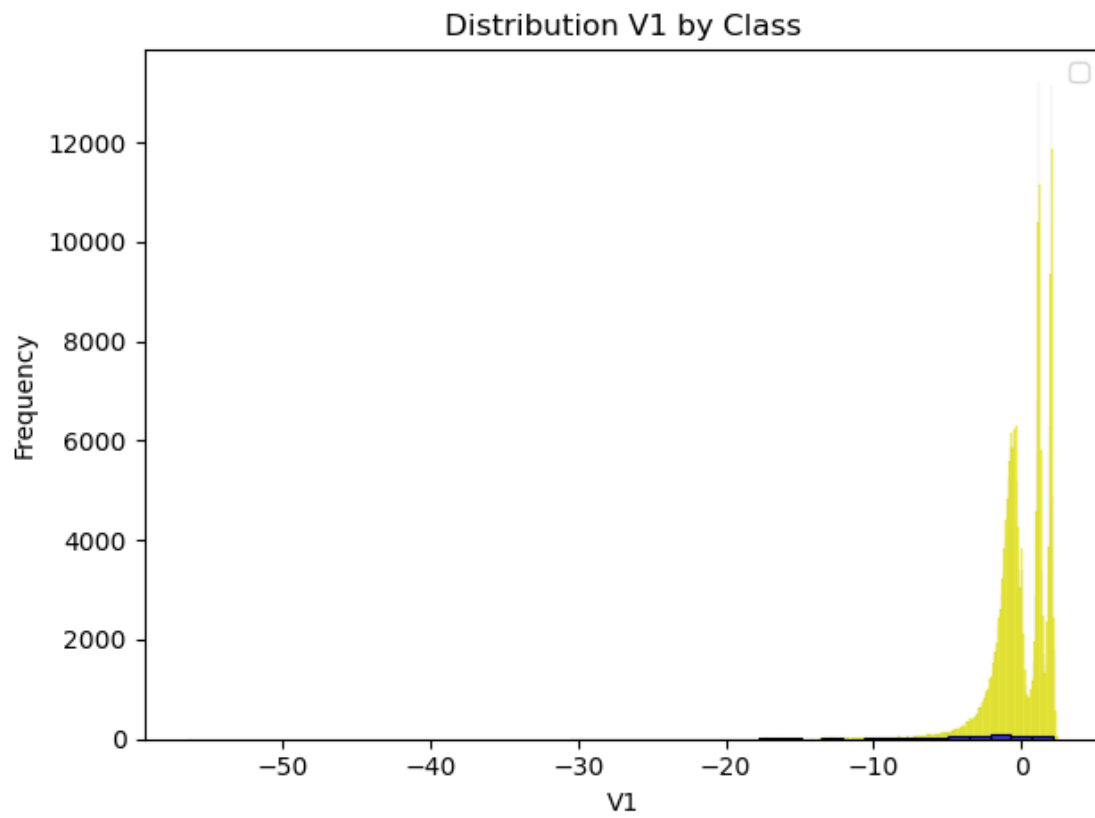
Weak Correlations: Most of the other features exhibit weak correlations with each other, which is common in high-dimensional datasets. This suggests that the features are relatively independent of one another.

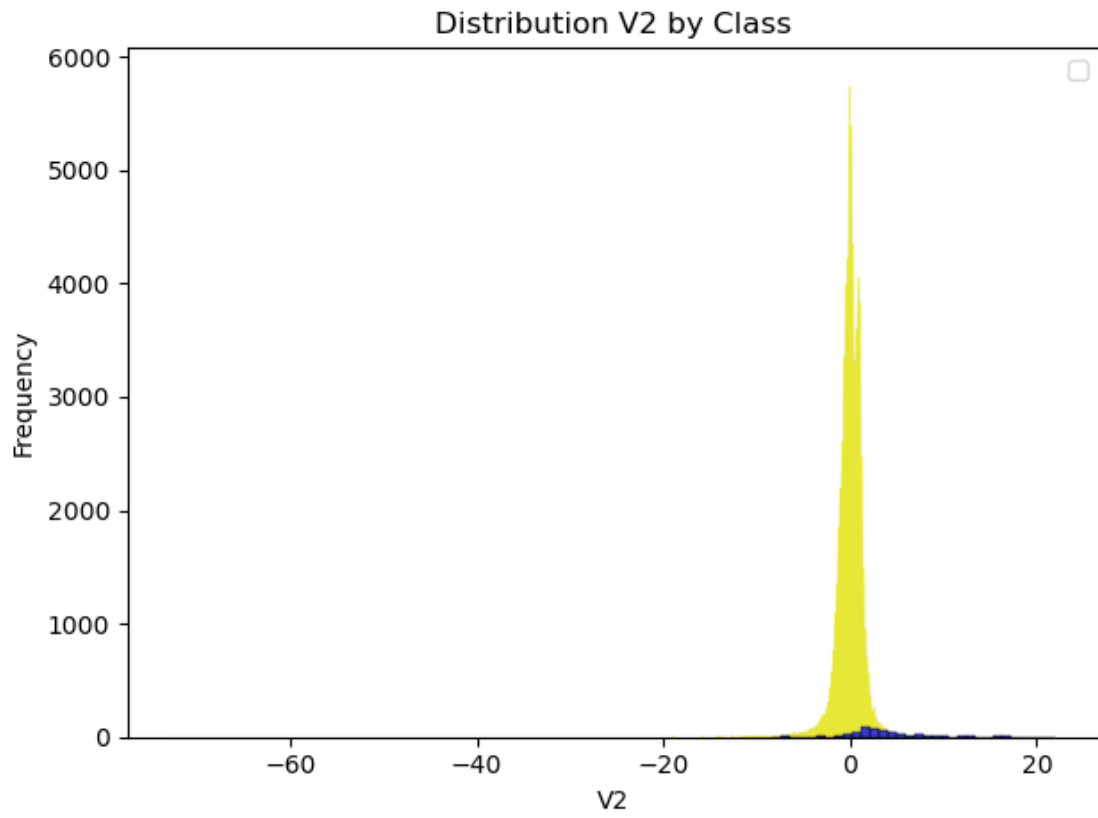
Target Variable (Class): The target variable (Class) shows a weak negative correlation with several features, indicating that there is no strong linear relationship between the features and the likelihood of fraud. This is typical in fraud detection datasets, where the relationship may be more complex.

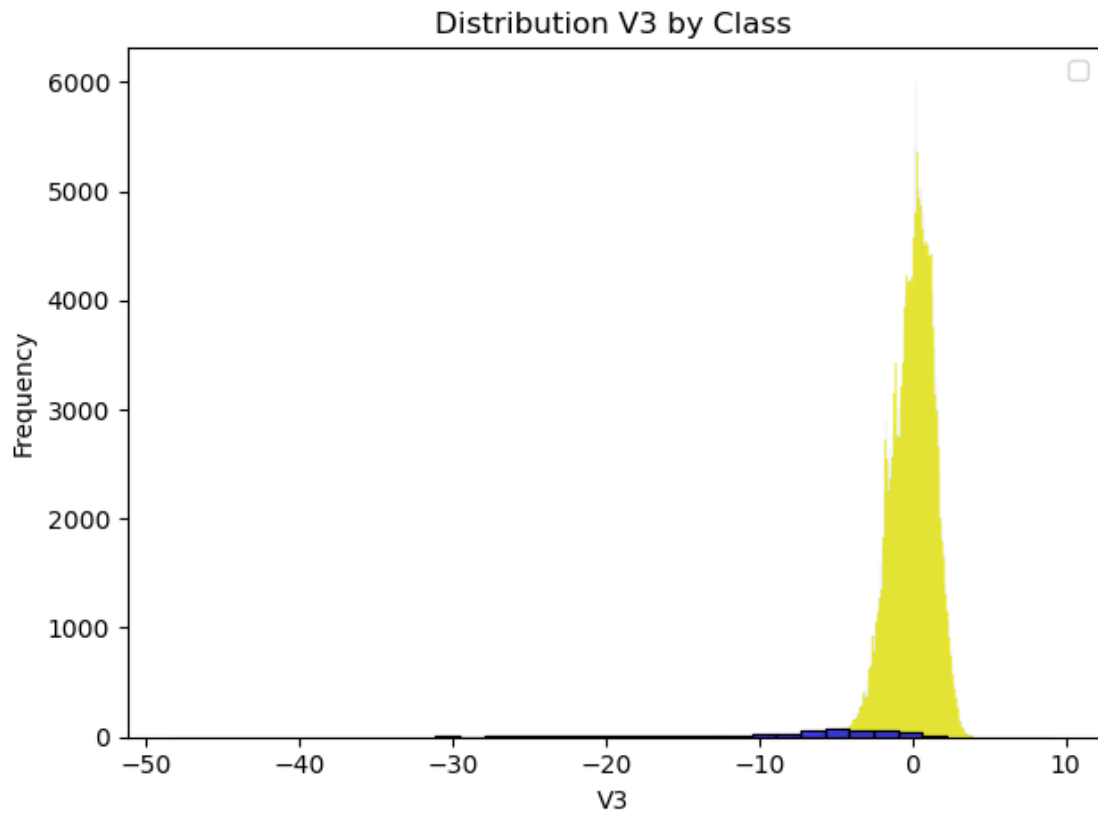
```
[9]: # Compare feature vs. target columns
for i in A:
    sns.histplot(x=A[i][data['Class'] == 0], color='yellow')
    sns.histplot(x=A[i][data['Class'] == 1], color='blue')

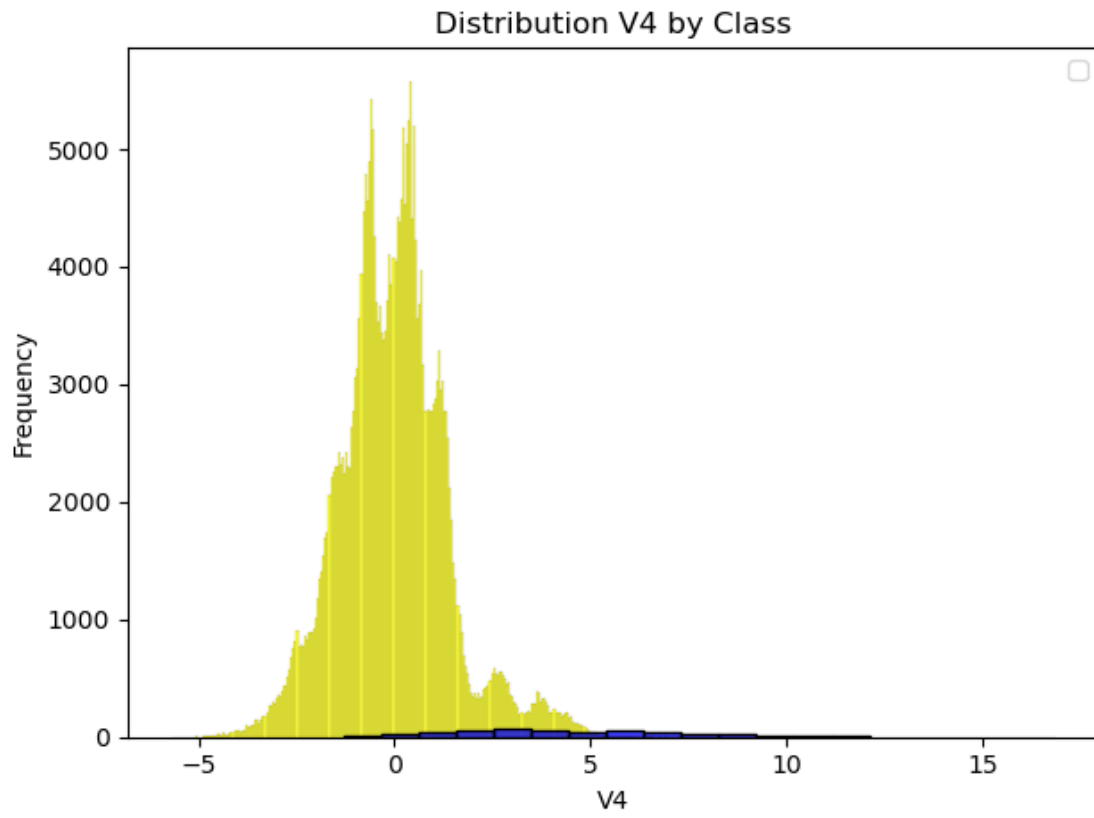
    plt.title(f'Distribution {i} by Class')
    plt.xlabel(i)
    plt.ylabel('Frequency')
    plt.legend()
    plt.tight_layout()
    plt.show()
```

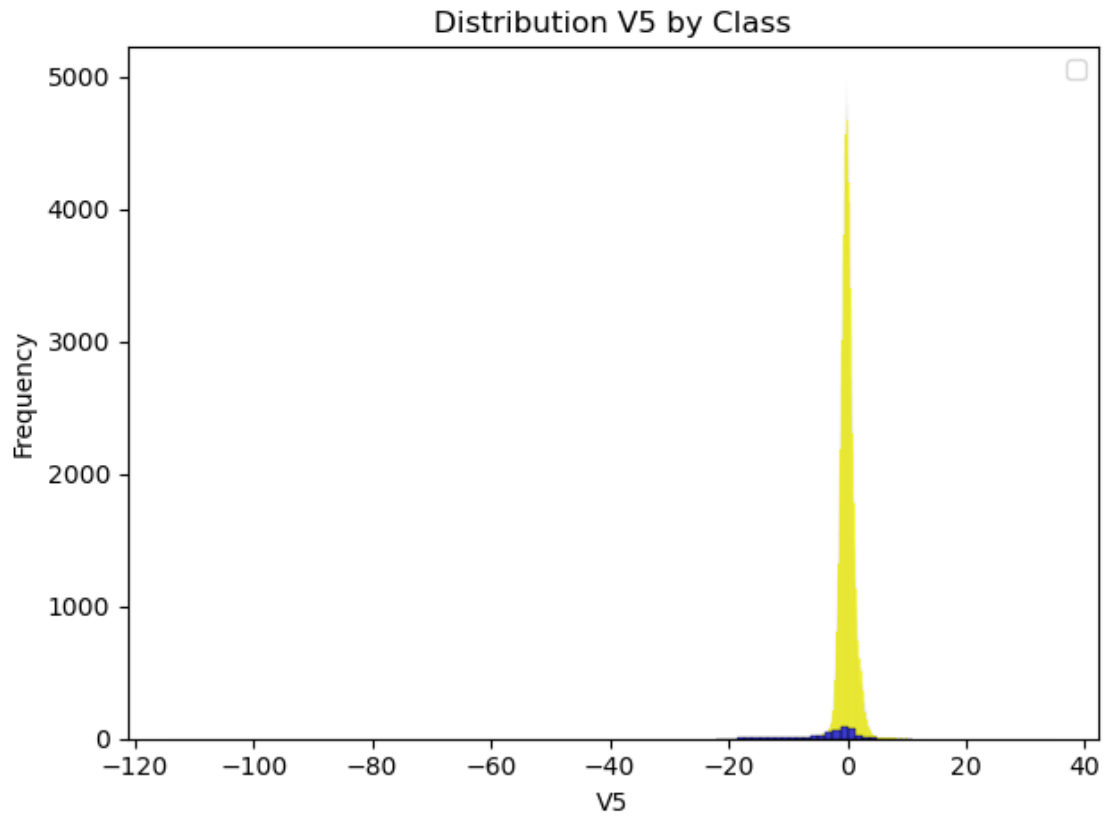


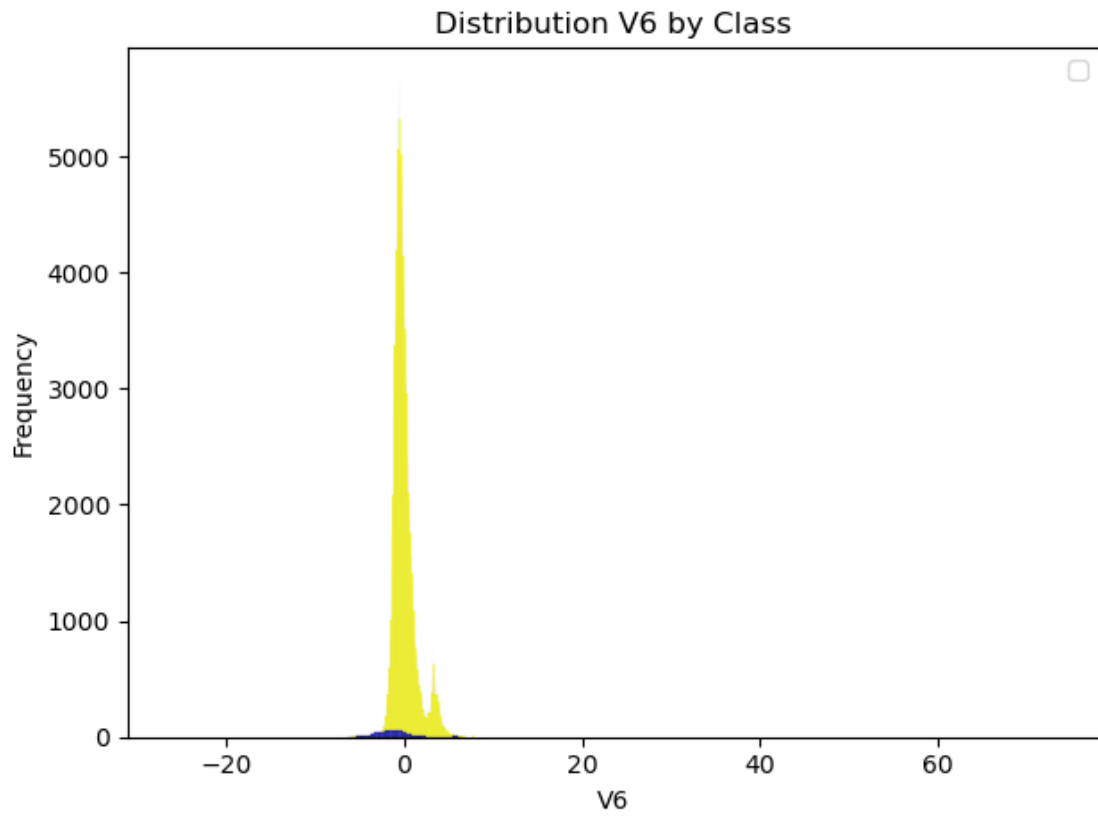


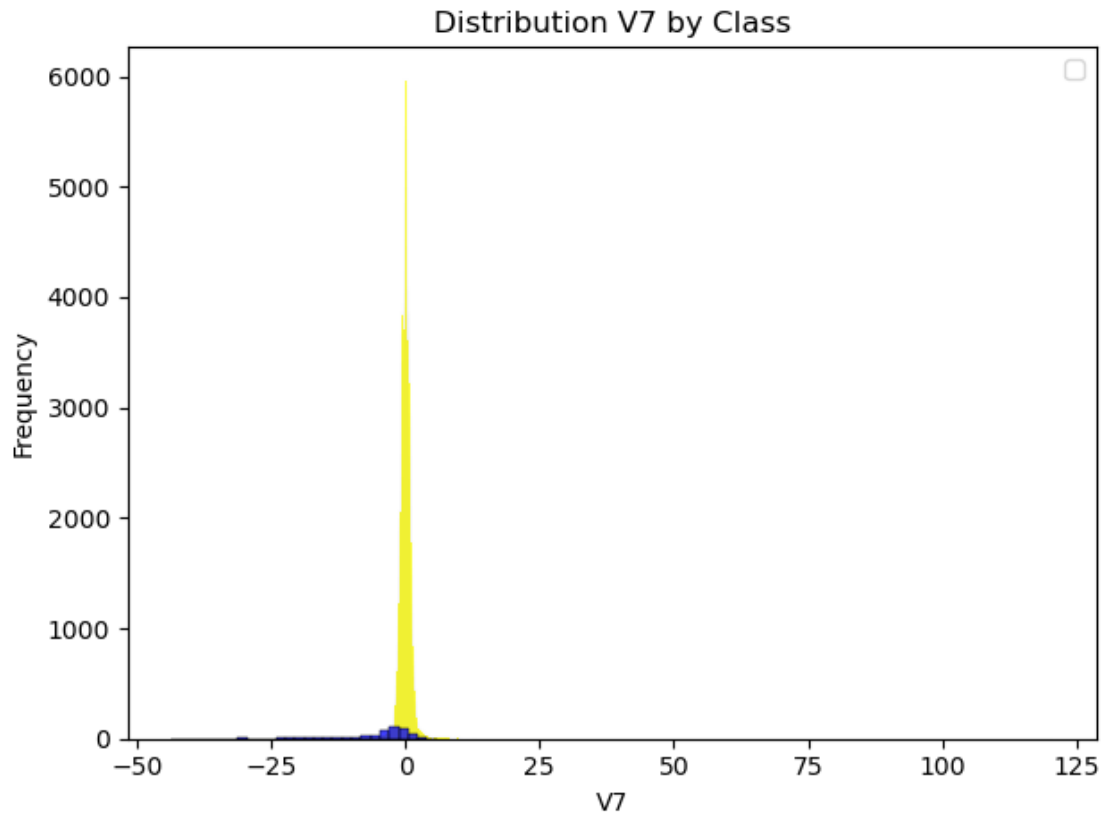


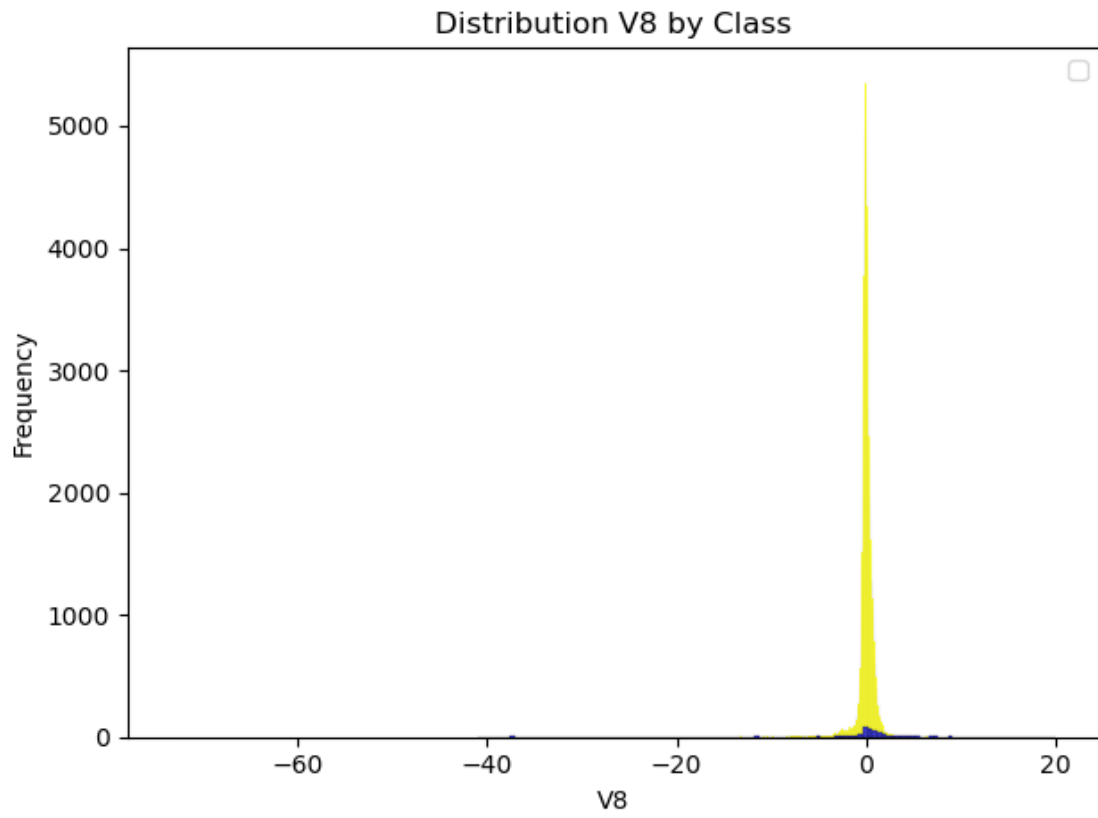


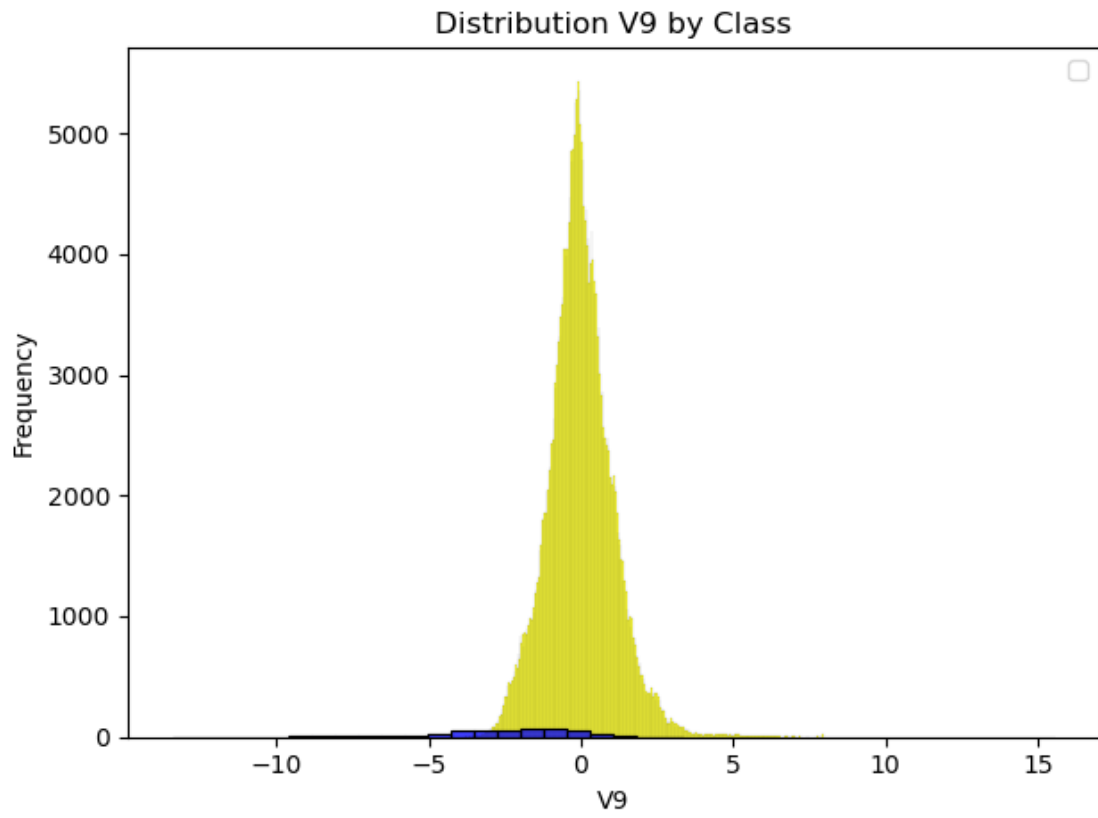


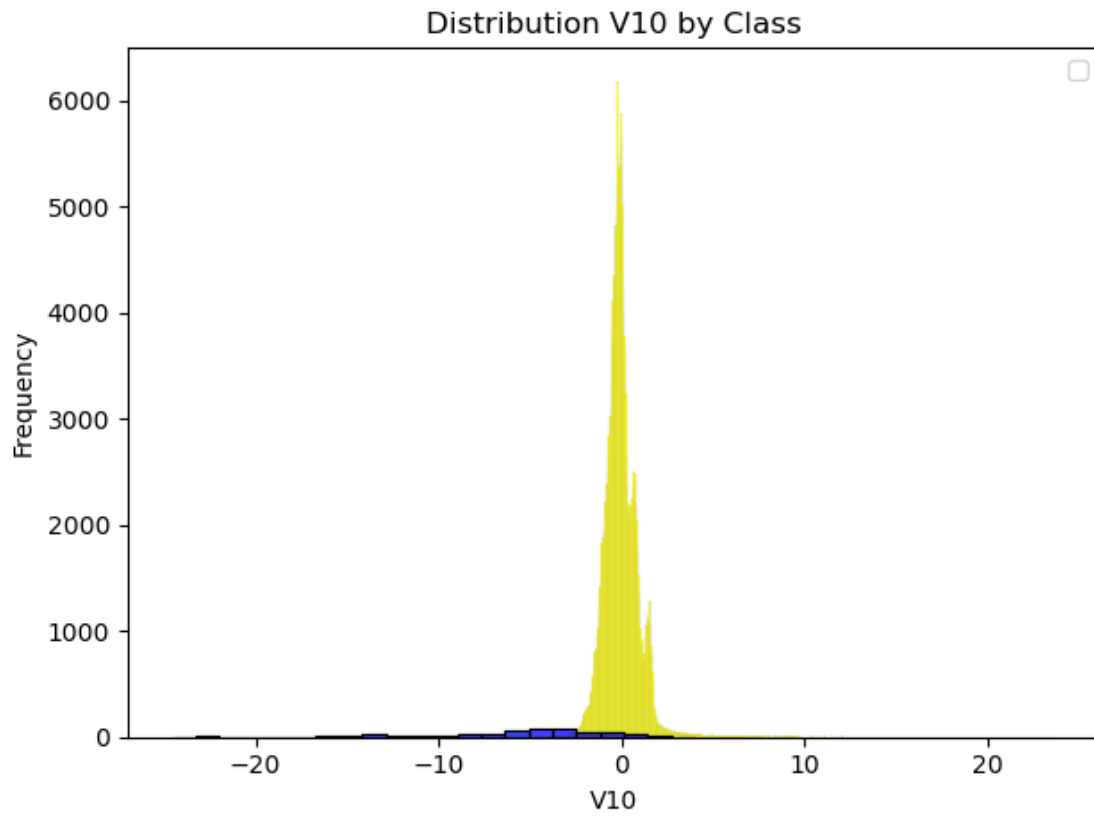


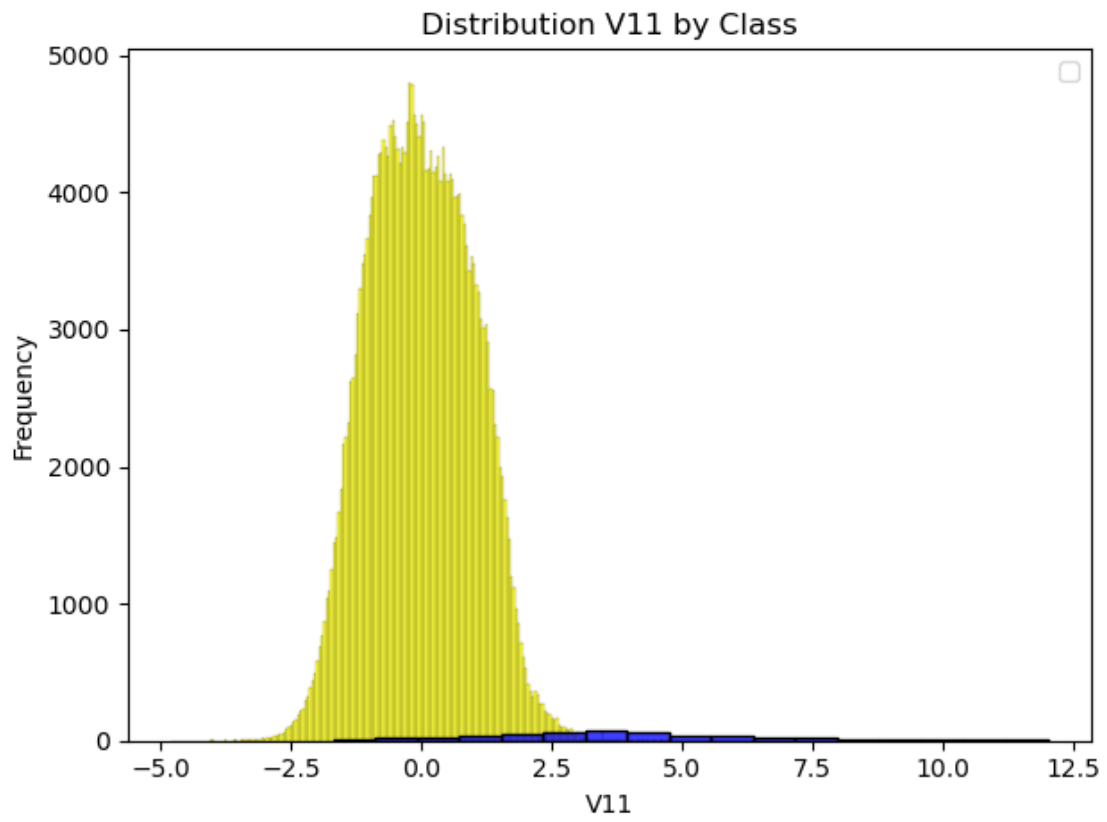


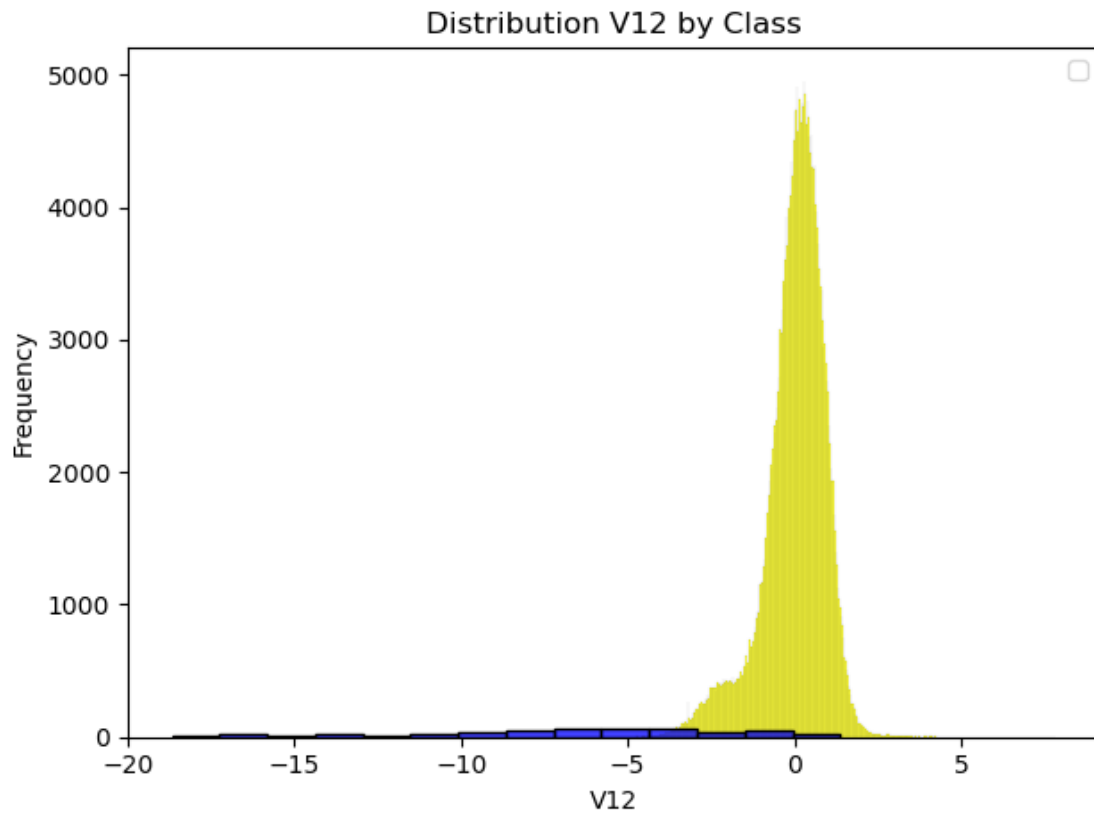


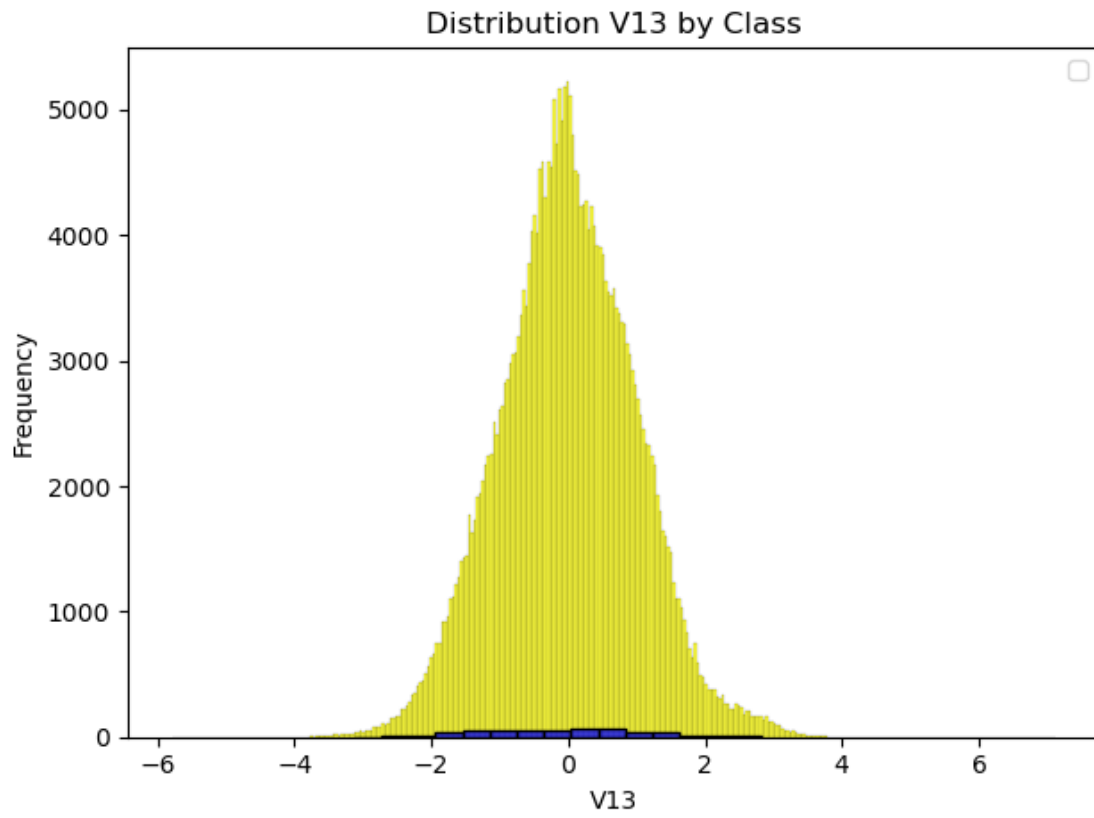


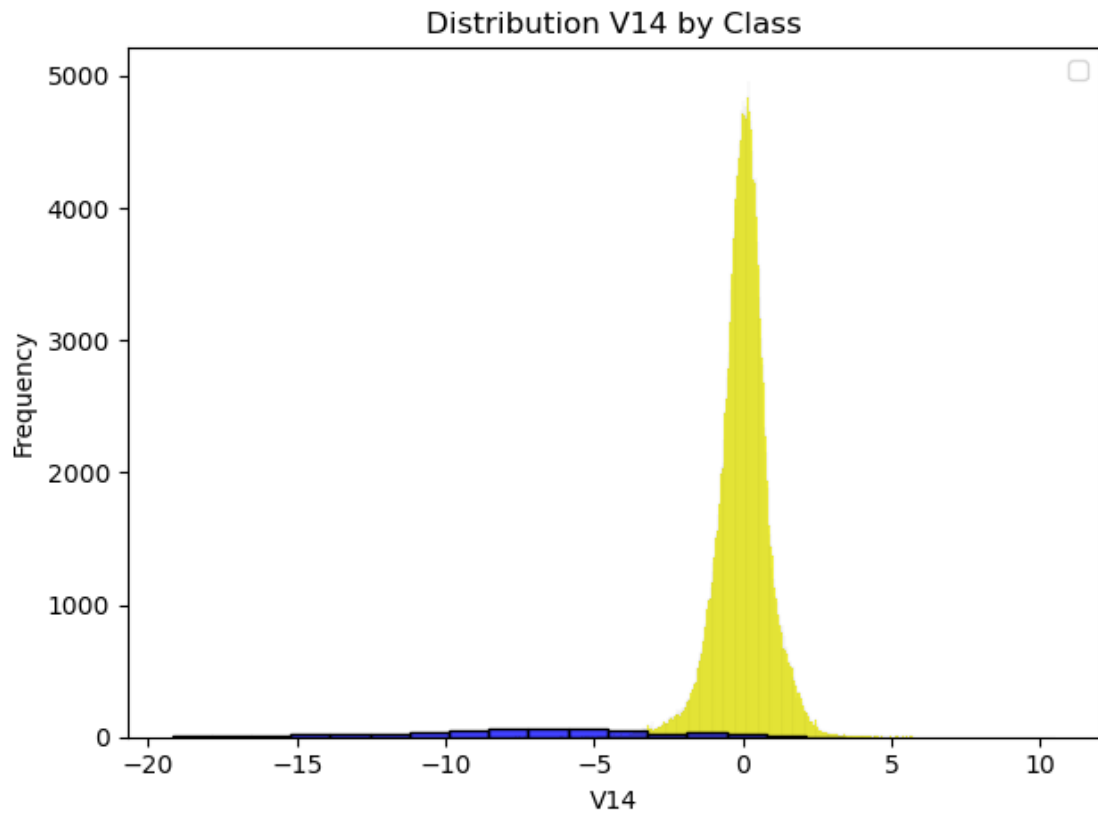


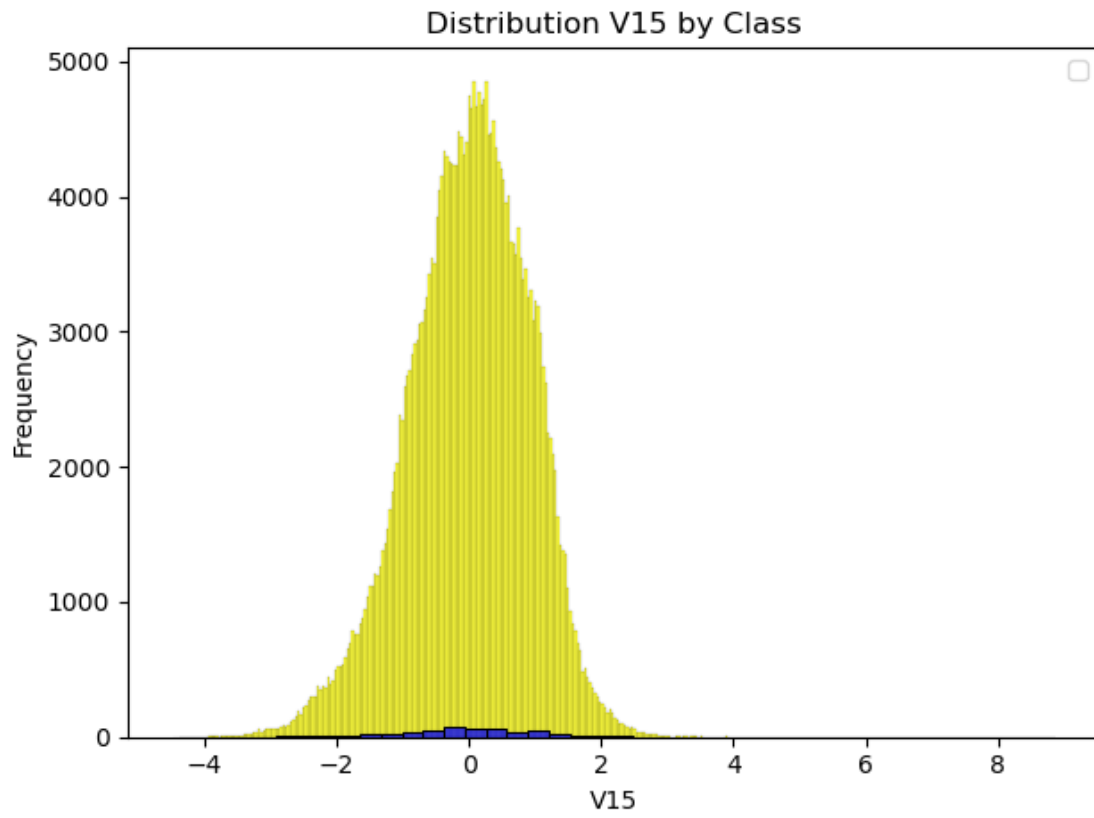


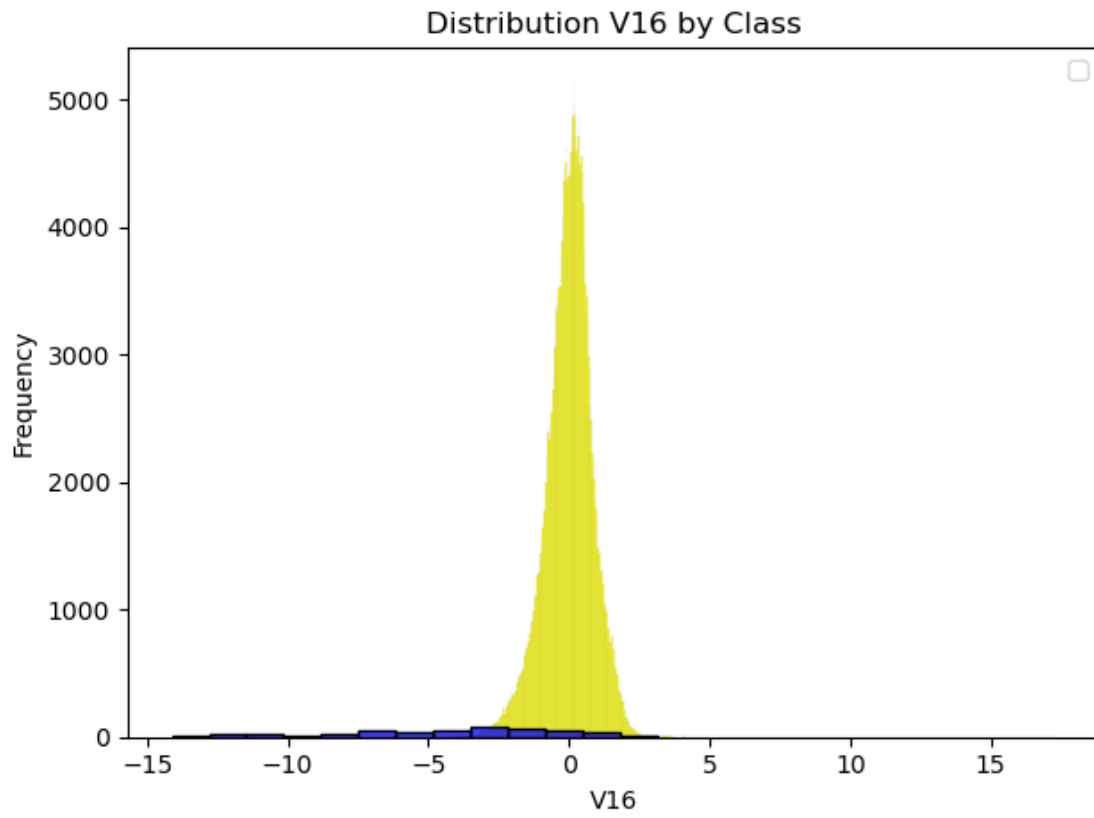


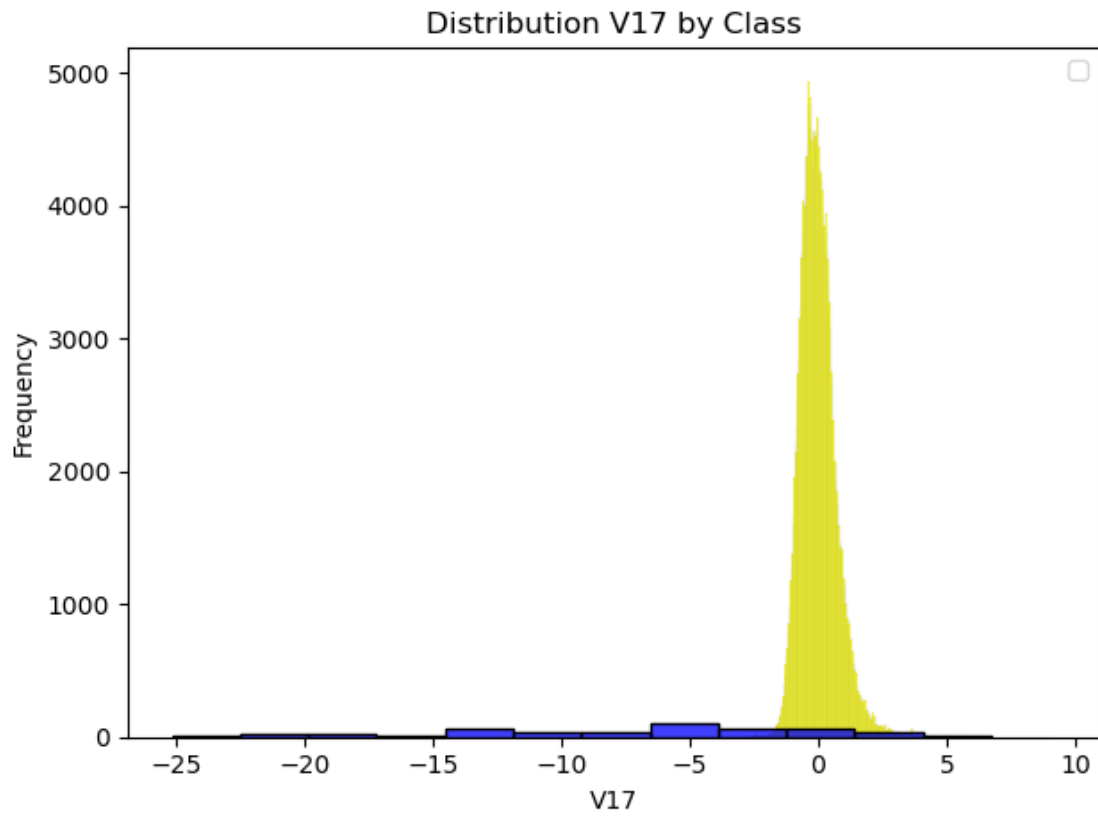


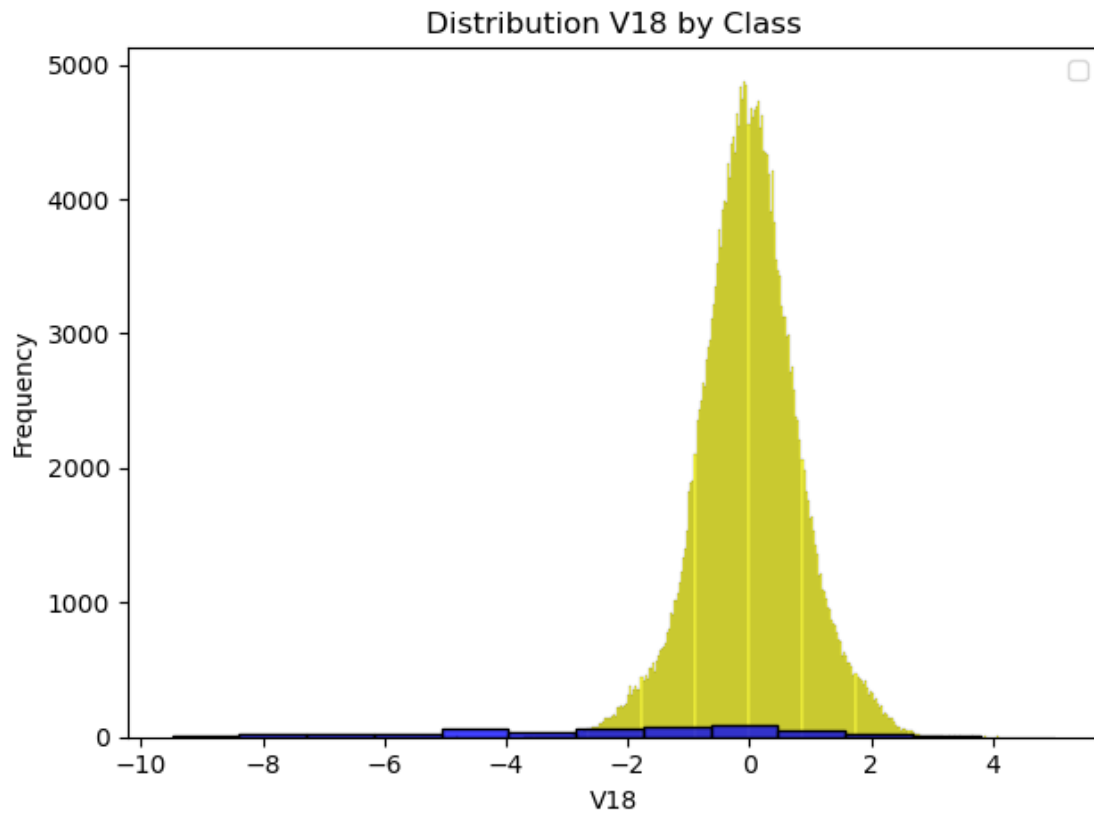


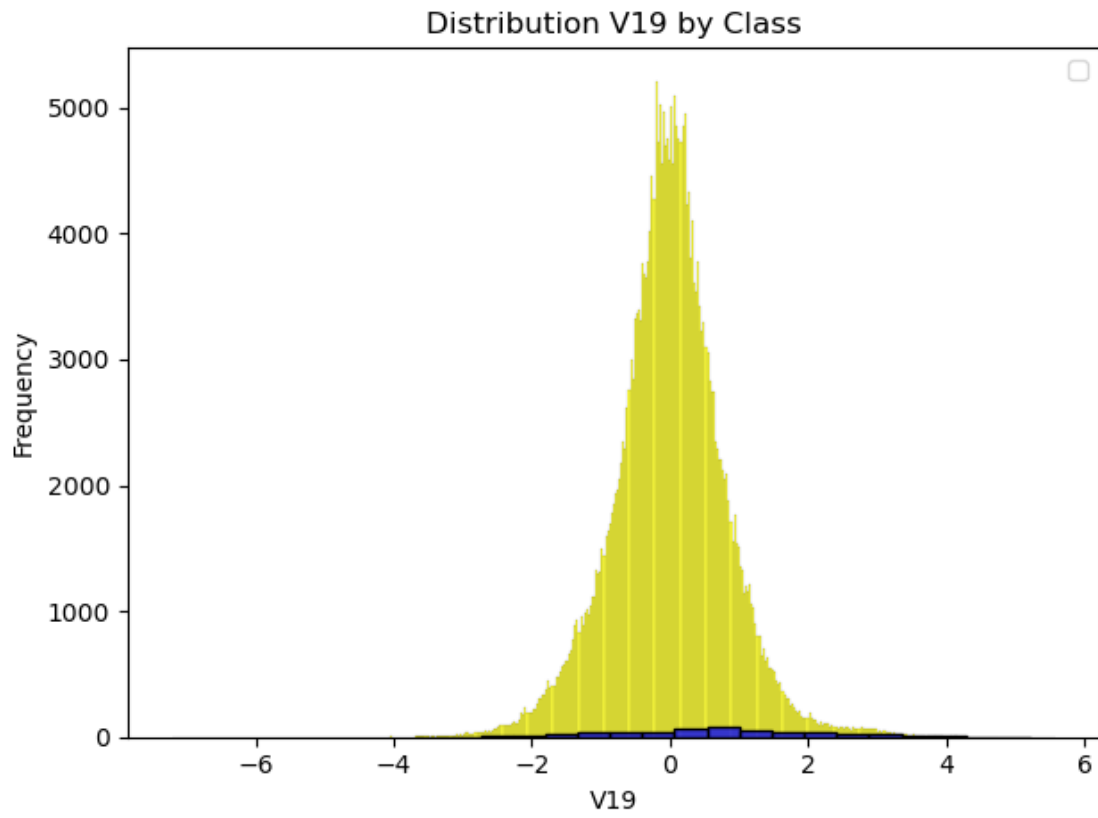


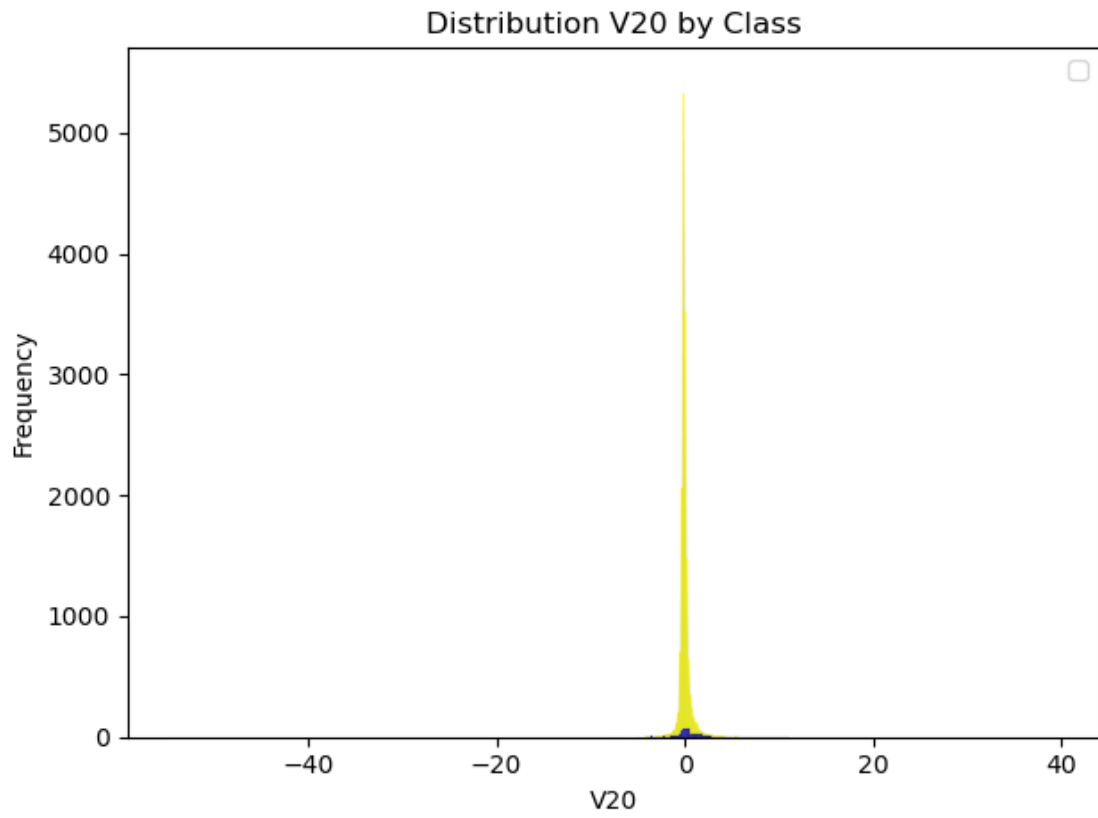


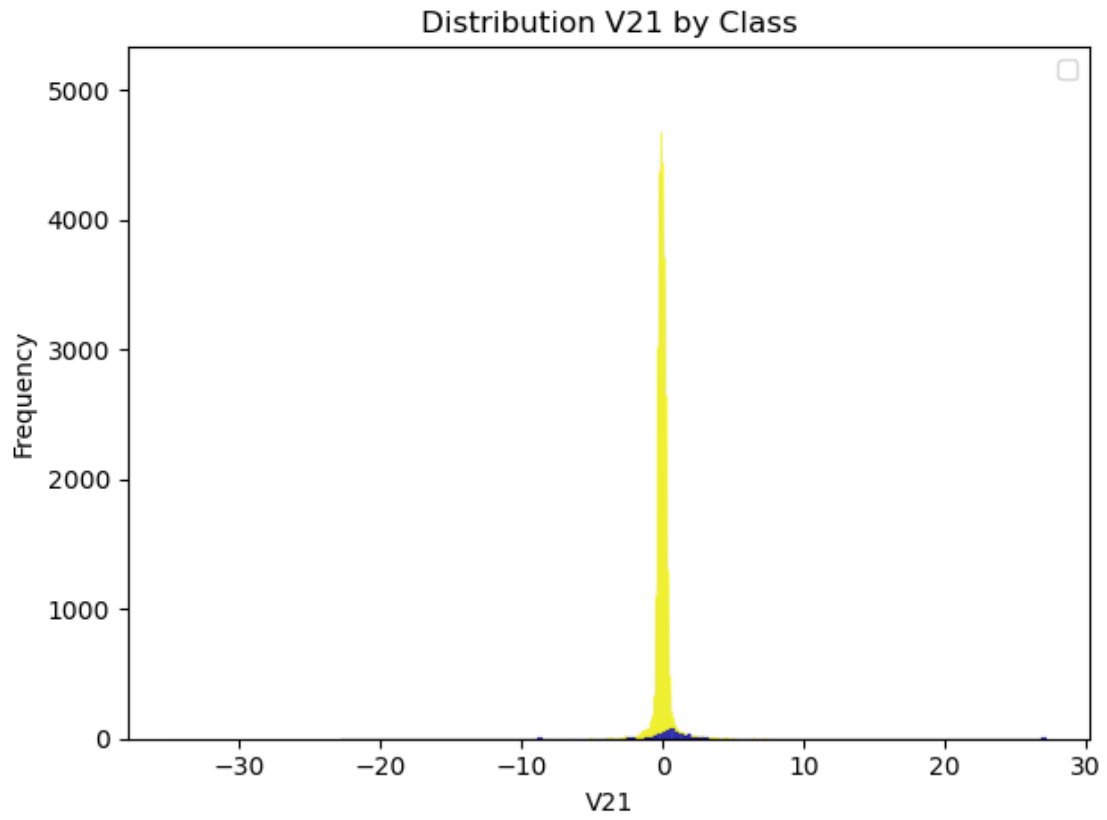


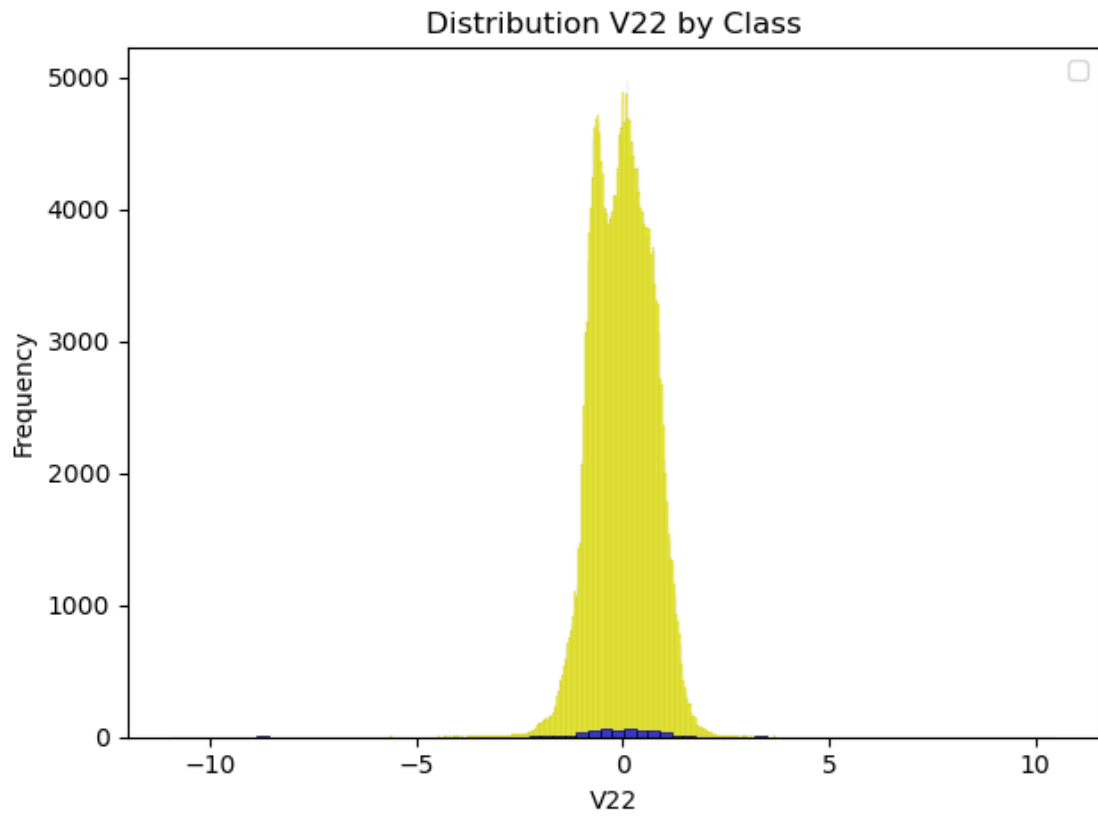


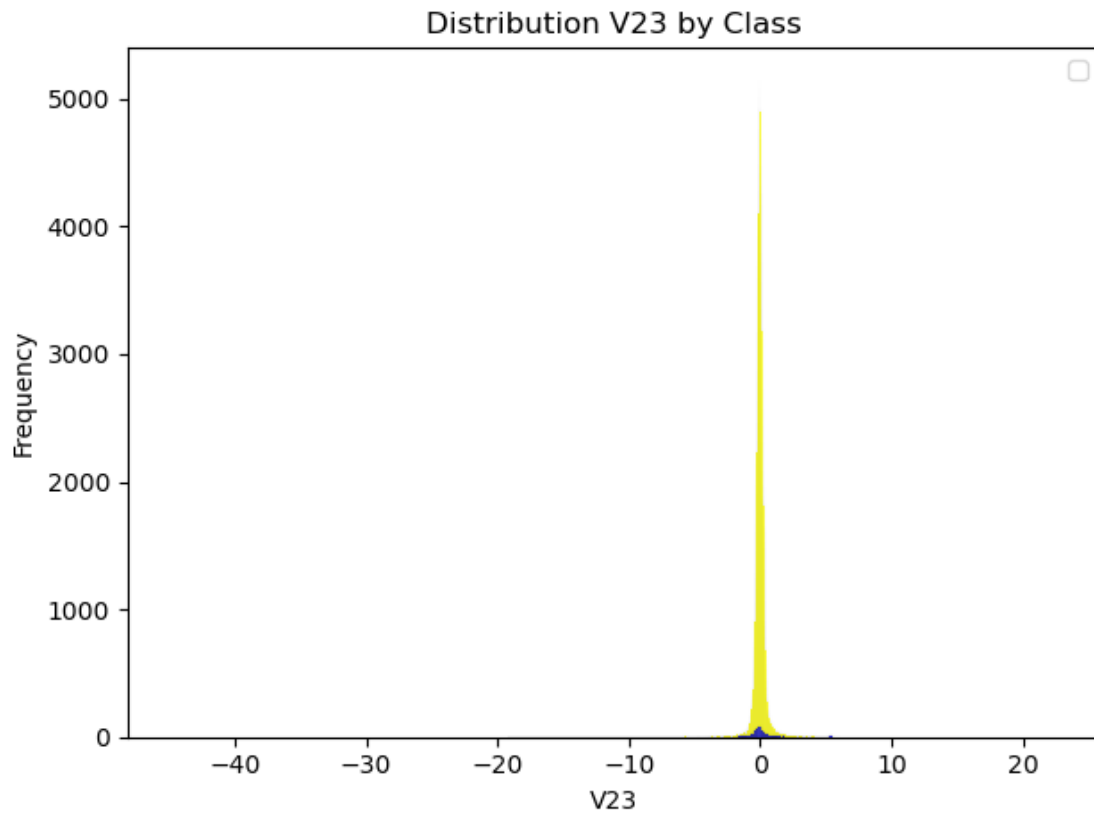


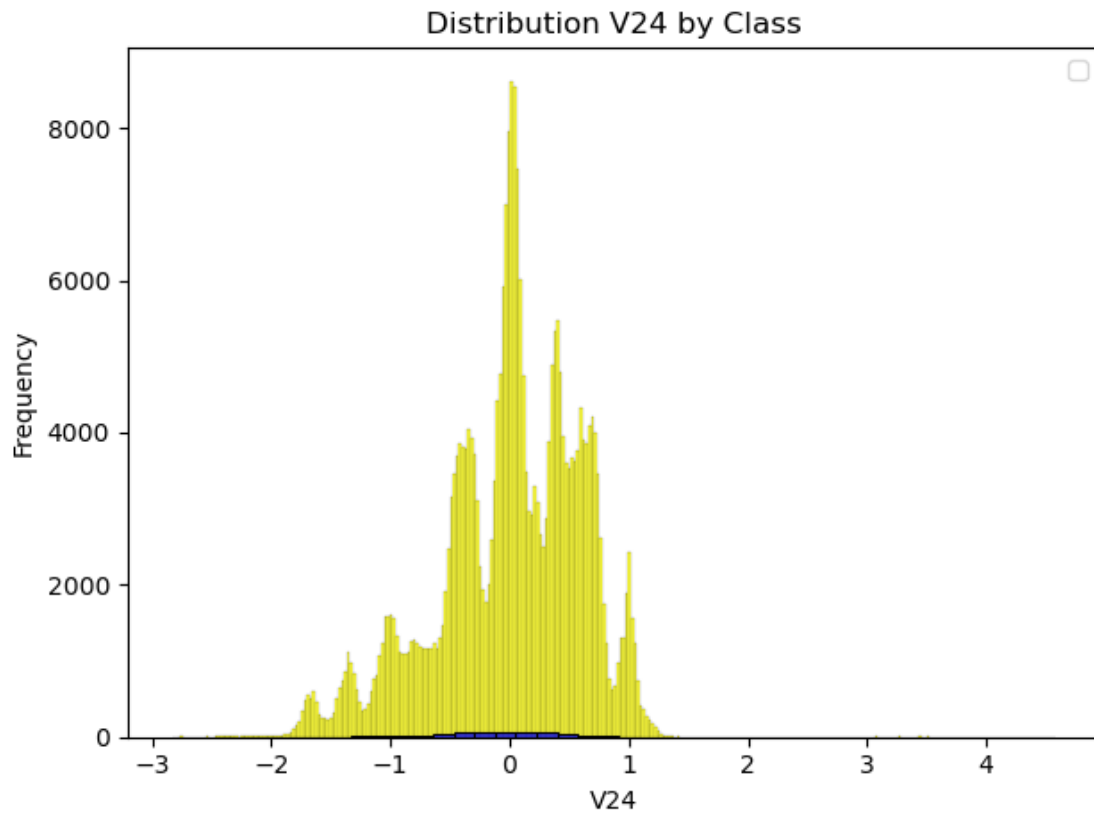


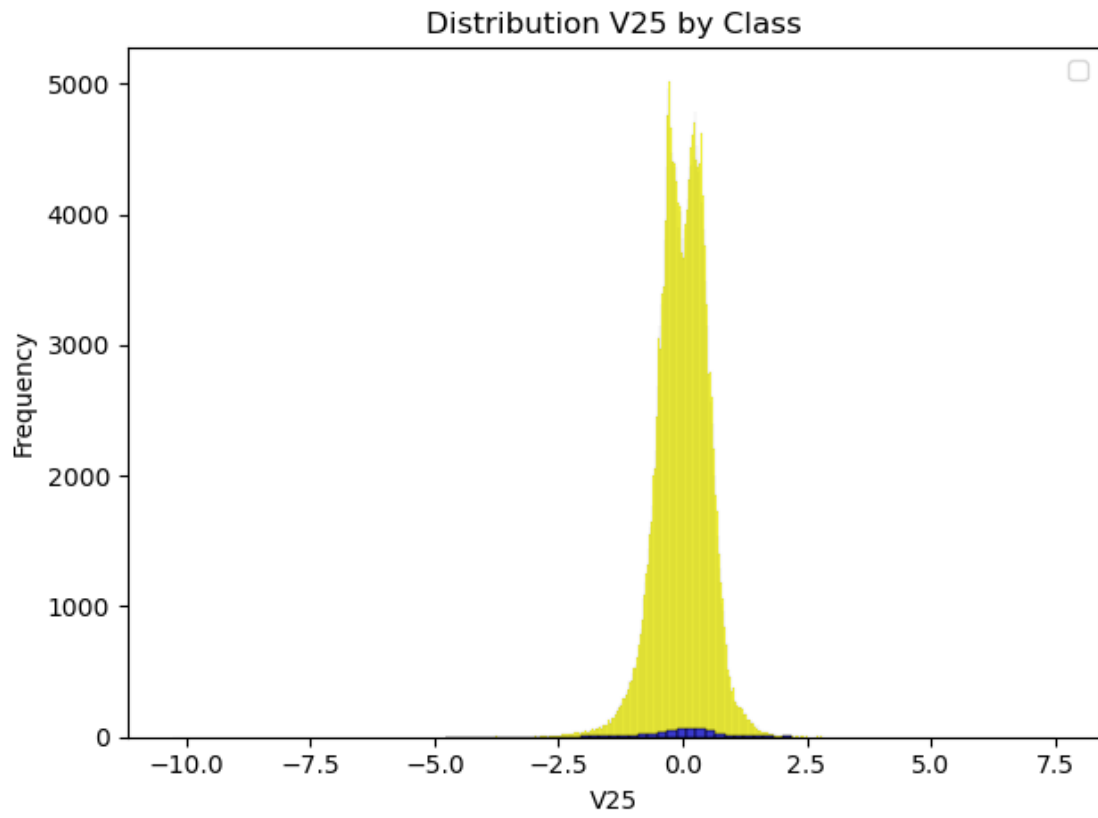


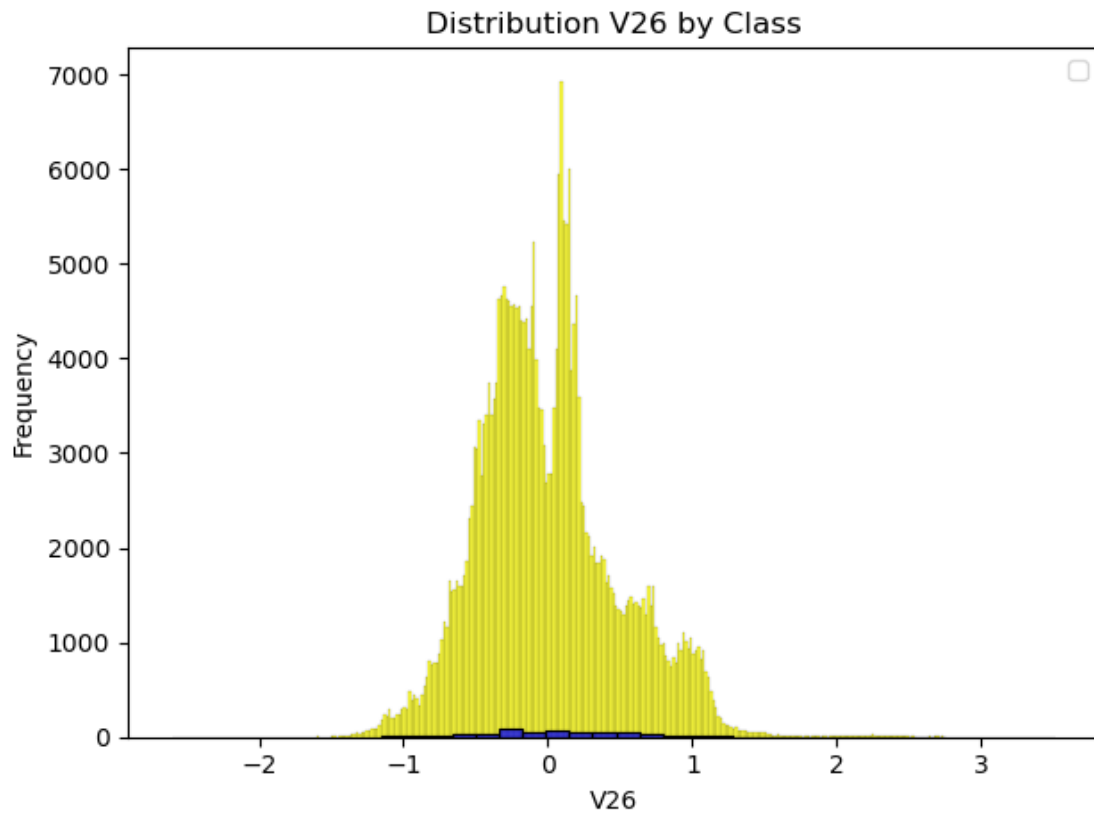


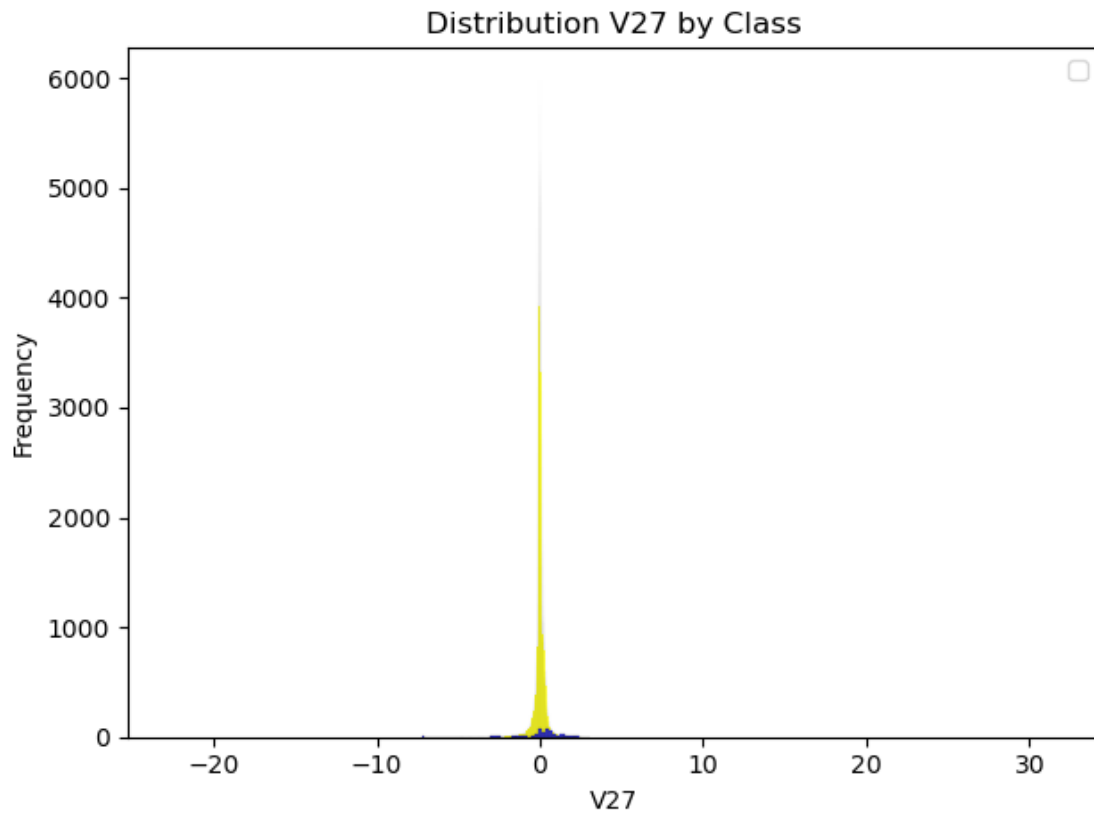


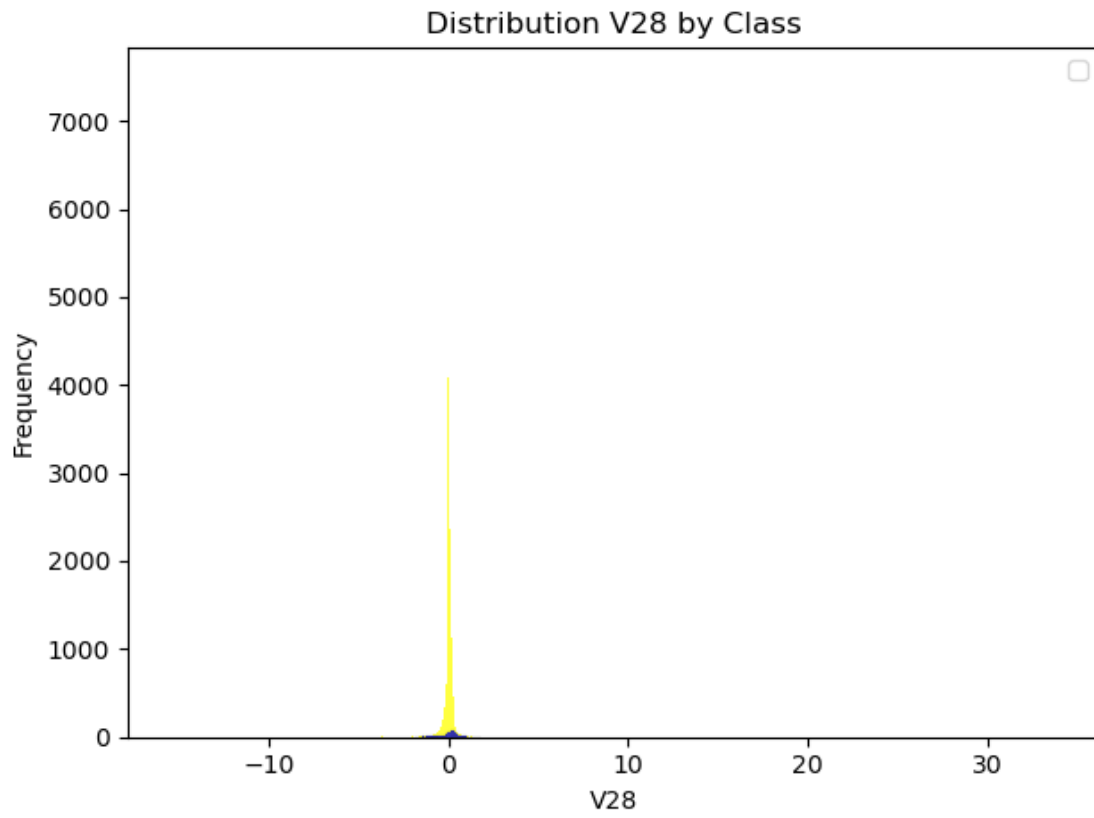


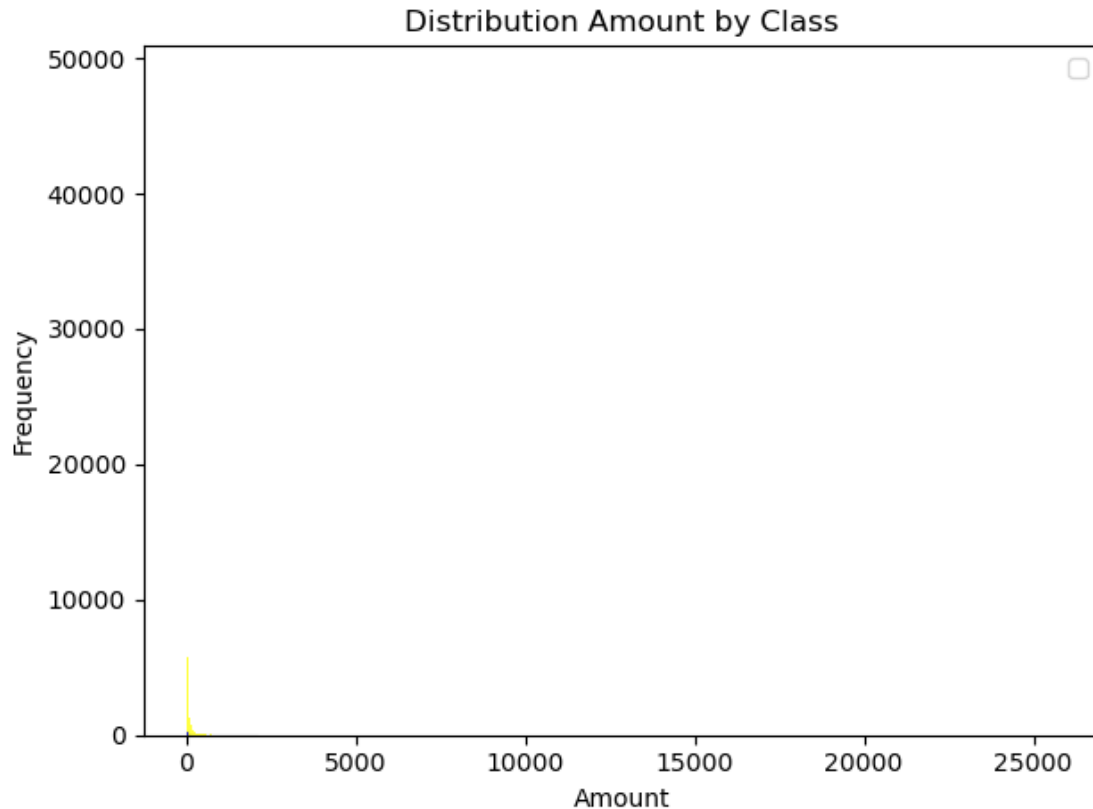












Insights of correlation matrix and histogram distributions of each feature for Class :

1 **Time:** The distribution is similar for both classes, indicating that the time of transaction is not a strong indicator of fraud.

2 **V1 to V28:** These features are PCA components, so their distributions are not directly interpretable. However, some features show distinct differences:

- **V4, V11, V17:** These features have noticeable differences between the two classes, suggesting they may be important for distinguishing fraud.
 - **V10, V12, V14, V16:** These features also show some separation, indicating potential usefulness in fraud detection.
- 3 **Amount:** Fraudulent transactions tend to have lower amounts compared to non-fraudulent ones, which might be useful for prediction.

Overall, features with distinct differences in distribution between the two classes are likely more informative for predicting fraudulent transactions

Data Preprocessing

```
[12]: from sklearn.preprocessing import StandardScaler
#Scale 'Amount' and 'Time' (other features are already standardized)
scaler = StandardScaler()
A[['Time', 'Amount']] = scaler.fit_transform(A[['Time', 'Amount']])
```

```
# A is now scaled and ready for splitting
print(A.head())
```

	Time	V1	V2	V3	V4	V5	V6	\
0	-1.996583	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	
1	-1.996583	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	
2	-1.996562	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	
3	-1.996562	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	
4	-1.996541	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	

	V7	V8	V9	...	V20	V21	V22	V23	\
0	0.239599	0.098698	0.363787	...	0.251412	-0.018307	0.277838	-0.110474	
1	-0.078803	0.085102	-0.255425	...	-0.069083	-0.225775	-0.638672	0.101288	
2	0.791461	0.247676	-1.514654	...	0.524980	0.247998	0.771679	0.909412	
3	0.237609	0.377436	-1.387024	...	-0.208038	-0.108300	0.005274	-0.190321	
4	0.592941	-0.270533	0.817739	...	0.408542	-0.009431	0.798278	-0.137458	

	V24	V25	V26	V27	V28	Amount
0	0.066928	0.128539	-0.189115	0.133558	-0.021053	0.244964
1	-0.339846	0.167170	0.125895	-0.008983	0.014724	-0.342475
2	-0.689281	-0.327642	-0.139097	-0.055353	-0.059752	1.160686
3	-1.175575	0.647376	-0.221929	0.062723	0.061458	0.140534
4	0.141267	-0.206010	0.502292	0.219422	0.215153	-0.073403

[5 rows x 30 columns]

Train-Test Split

```
[14]: from sklearn.model_selection import train_test_split

# Split into training and testing sets (80/20)
X_train, X_test, y_train, y_test = train_test_split(
    A, B, test_size=0.2, random_state=42, stratify=B) # maintain class
↪distribution
```

Handle Class Imbalance with SMOTE Fraud cases are extremely rare, so let's balance the training set using SMOTE (Synthetic Minority Over-sampling Technique):

```
[19]: from imblearn.over_sampling import SMOTE

# Apply SMOTE to the training data
smote = SMOTE(random_state=42)
X_train_resampled, y_train_resampled = smote.fit_resample(X_train, y_train)

# Check class distribution
from collections import Counter
print("Before SMOTE:", Counter(y_train))
print("After SMOTE:", Counter(y_train_resampled))
```

Before SMOTE: Counter({0: 227451, 1: 394})
After SMOTE: Counter({0: 227451, 1: 227451})

Evaluate multiple Classification models to compare performance

Model 1 - Logistic Regression

```
[21]: from sklearn.linear_model import LogisticRegression
      from sklearn.metrics import classification_report, confusion_matrix, \
      roc_auc_score

      # Initialize and train the model
      model = LogisticRegression(max_iter=1000, random_state=42)
      model.fit(X_train_resampled, y_train_resampled)

      # Predict on the test set
      y_pred = model.predict(X_test)
      y_proba = model.predict_proba(X_test)[:, 1]

[23]: # Classification report
      print("Classification Report:\n", classification_report(y_test, y_pred))

      # Confusion matrix
      print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))

      # ROC AUC score
      print("ROC AUC Score:", roc_auc_score(y_test, y_proba))
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.97	0.99	56864
1	0.06	0.92	0.11	98
accuracy			0.97	56962
macro avg	0.53	0.95	0.55	56962
weighted avg	1.00	0.97	0.99	56962

Confusion Matrix:

```
[[55406 1458]
 [    8   90]]
```

ROC AUC Score: 0.9698482164390798

Model 2 - Random Forest

```
[34]: from sklearn.ensemble import RandomForestClassifier

      rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
      rf_model.fit(X_train_resampled, y_train_resampled)
```

```

rf_pred = rf_model.predict(X_test)
rf_proba = rf_model.predict_proba(X_test)[:, 1]

print("=== Random Forest ===")
print("Confusion Matrix:\n", confusion_matrix(y_test, rf_pred))
print("Classification Report:\n", classification_report(y_test, rf_pred))
print("ROC AUC Score:", roc_auc_score(y_test, rf_proba))

```

```

=== Random Forest ===
Confusion Matrix:
[[56849   15]
 [   16   82]]
Classification Report:

```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56864
1	0.85	0.84	0.84	98
accuracy			1.00	56962
macro avg	0.92	0.92	0.92	56962
weighted avg	1.00	1.00	1.00	56962

ROC AUC Score: 0.9731024901519414

Model 3 - K-Nearest Neighbors

```

[27]: from sklearn.neighbors import KNeighborsClassifier
knn_model = KNeighborsClassifier(n_neighbors=5)
knn_model.fit(X_train_resampled, y_train_resampled)

knn_pred = knn_model.predict(X_test)
knn_proba = knn_model.predict_proba(X_test)[:, 1]

print("\n=== K-Nearest Neighbors ===")
print("Confusion Matrix:\n", confusion_matrix(y_test, knn_pred))
print("Classification Report:\n", classification_report(y_test, knn_pred))
print("ROC AUC Score:", roc_auc_score(y_test, knn_proba))

```

```

=== K-Nearest Neighbors ===
Confusion Matrix:
[[56765   99]
 [   12   86]]
Classification Report:

```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	56864
1	0.46	0.88	0.61	98

accuracy			1.00	56962
macro avg	0.73	0.94	0.80	56962
weighted avg	1.00	1.00	1.00	56962

ROC AUC Score: 0.9535882427675628

Visualization of ROC curves

```
[36]: from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt

# we have already calculated Predict probabilities for all models above

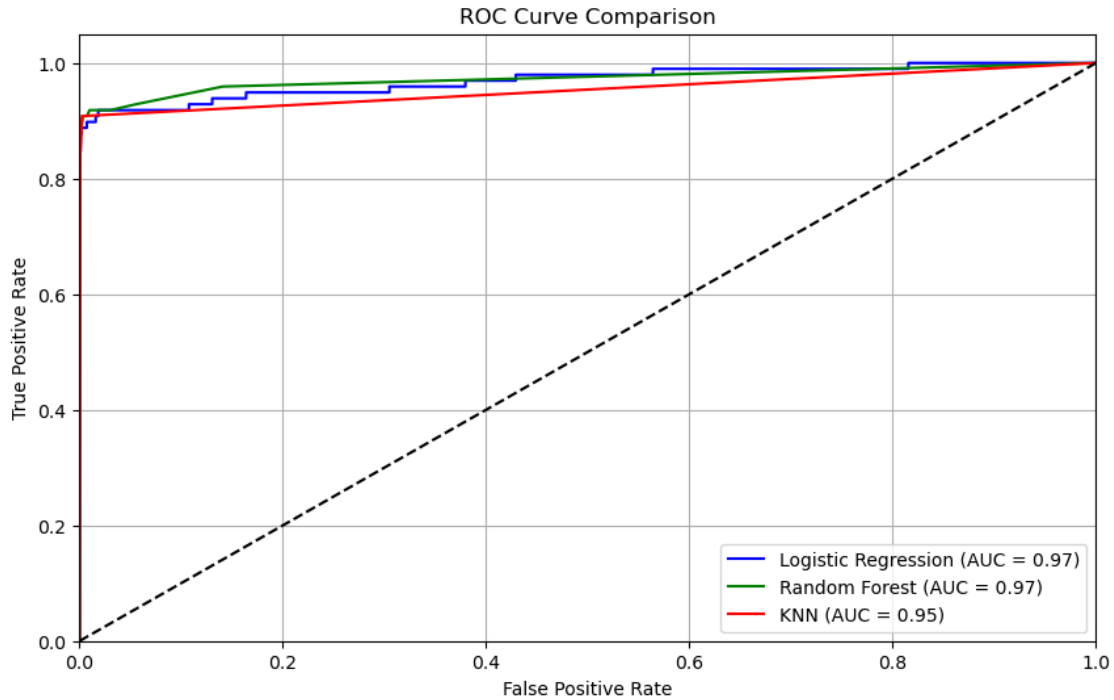
# Compute ROC curve and AUC
fpr_lr, tpr_lr, _ = roc_curve(y_test, y_proba)
roc_auc_lr = auc(fpr_lr, tpr_lr)

fpr_rf, tpr_rf, _ = roc_curve(y_test, rf_proba)
roc_auc_rf = auc(fpr_rf, tpr_rf)

fpr_knn, tpr_knn, _ = roc_curve(y_test, knn_proba)
roc_auc_knn = auc(fpr_knn, tpr_knn)

# Plot all ROC curves
plt.figure(figsize=(10, 6))
plt.plot(fpr_lr, tpr_lr, label=f'Logistic Regression (AUC = {roc_auc_lr:.2f})',
         color='blue')
plt.plot(fpr_rf, tpr_rf, label=f'Random Forest (AUC = {roc_auc_rf:.2f})',
         color='green')
plt.plot(fpr_knn, tpr_knn, label=f'KNN (AUC = {roc_auc_knn:.2f})', color='red')

plt.plot([0, 1], [0, 1], 'k--') # Diagonal reference line
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve Comparison')
plt.legend(loc="lower right")
plt.grid()
plt.show()
```



Model Performance Summary:

Logistic Regression - Accuracy= 0.97% , ROC AUC= 0.96

Random Forest - Accuracy= 1.00% , ROC AUC= 0.97

KNN - Accuracy= 1.00% , ROC AUC= 0.95

Conclusion

Based on the evaluation metrics and visual analysis, Random Forest outperformed other models in every critical category—achieving high accuracy, and ROC AUC. This indicates that it can reliably detect fraudulent transactions while minimizing false positives and false negatives. Logistic Regression also performed decently and can serve as a lightweight alternative. KNN, while simpler, lagged behind in key metrics, making it less suitable for highly imbalanced fraud detection tasks. Overall, this project demonstrates that combining robust preprocessing techniques with a powerful ensemble model like Random Forest is an effective strategy for fraud detection